

Deep Learning and Optimization

Unpacking Transformers, LLMs and Diffusion

Session 1

olivier.koch@ensae.fr

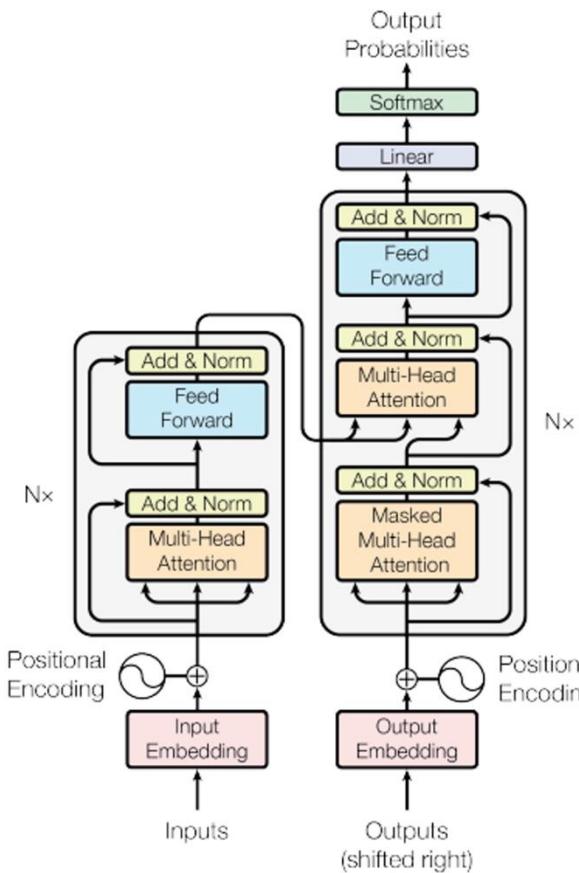
[slack #ensae-dl-2026](#)

What you will learn

- Concepts and practical aspects of deep learning
- Build your own backprop, mini-gpt and diffusion models

This course is not a complete overview of all ML & deep learning techniques.

Learning from scratch matters



mmm... so what?

Practical stuff

- 6 sessions
- 1.5h theory + 1.5h practice
- Grading: 6 notebooks (50%) + 1 quiz (50%)
- Each notebook should be yours
- I expect you to return your notebook at the end of each session.
- You can then send an updated version over the following week to aim for a better grade (and learn more).

On the usage of AI tools

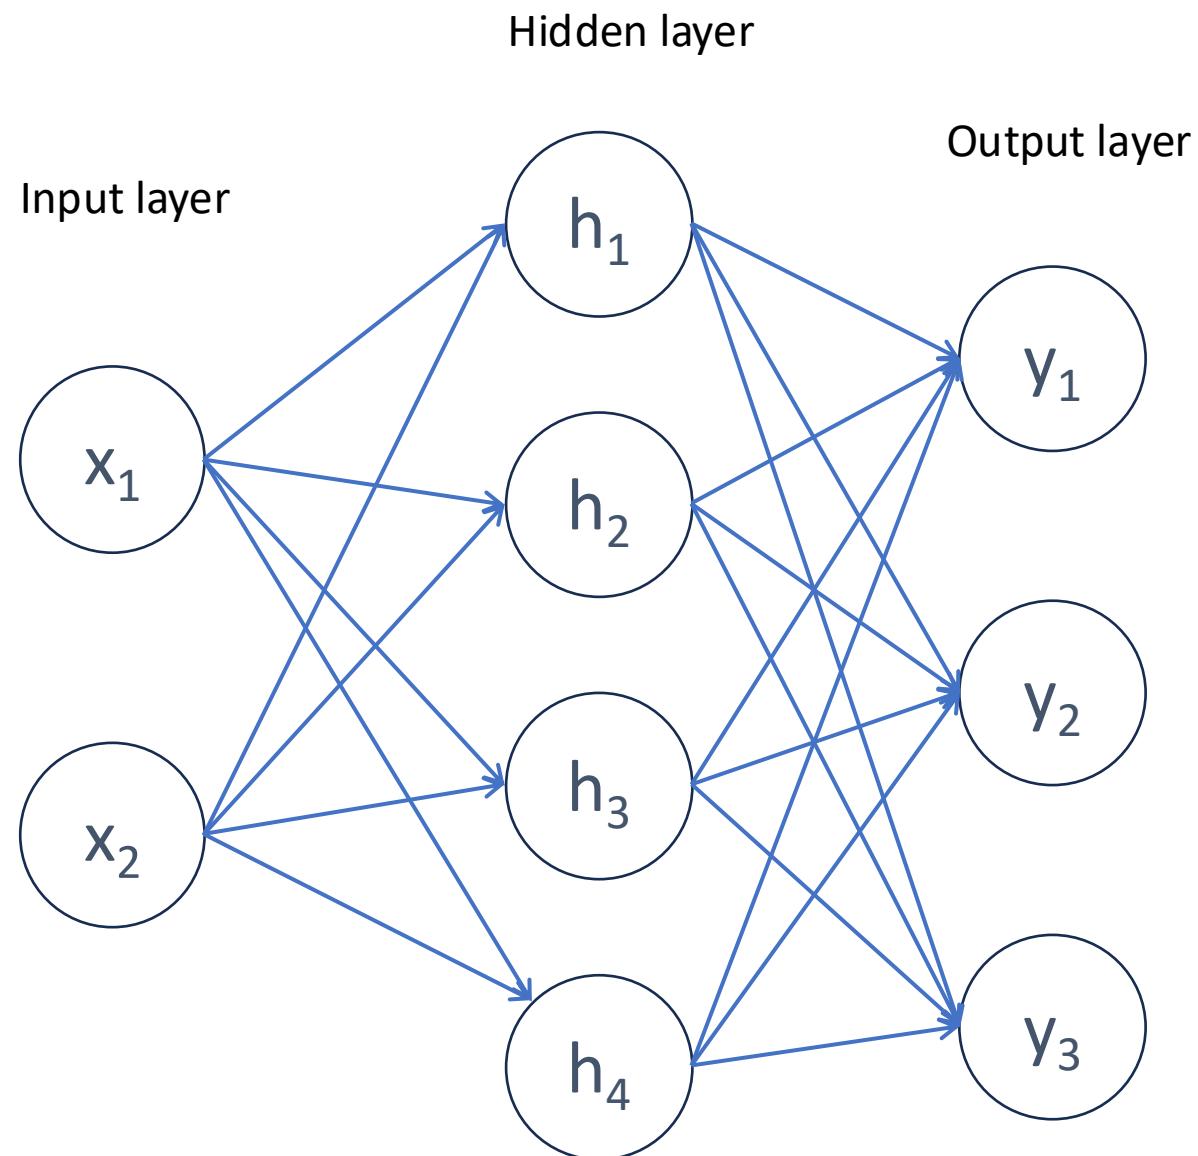
- This class will hardwire the fundamentals of AI in your brain
- You will only learn *by doing, from scratch*
- The use of AI tools in class is *discouraged*
- AI tools won't help you during the quiz (on paper, open questions)
- AI tools are the future!

References

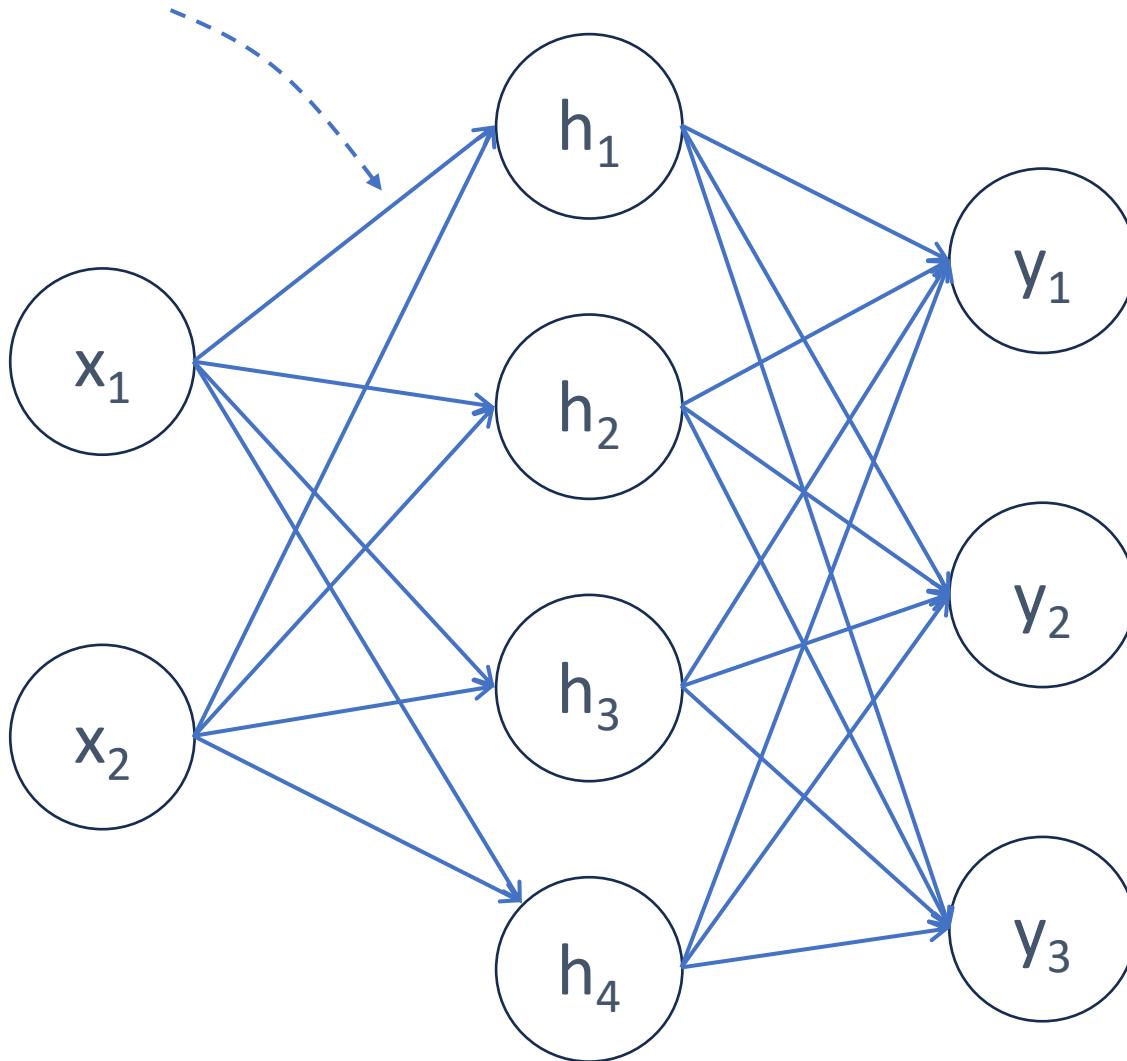
- [Understanding Deep Learning](#), Simon Prince, December 2023
- [Deep Learning: Foundations and Concepts](#), Christopher and Hugh Bishop, 2023
- [Neural networks: Zero to Hero](#), Andrej Karpathy, Youtube Lecture Series, 2022

| | Session | Date | Content |
|--------------|---------|---------|---|
| Foundations | 1 | Jan, 28 | Intro to DL TP: micrograd |
| | 2 | Feb, 4 | Fundamentals I: backprop, loss functions TP: bigram, MLP for next character prediction |
| | 3 | Feb, 11 | Fundamentals II: DL architectures TP: tensor-based models |
| Applications | 4 | Feb, 18 | Attention & Transformers TP: GPT from scratch |
| | 5 | Feb, 25 | DL for Computer vision TP: convnets on CIFAR-10 |
| | 6 | Mar, 3 | VAE and Diffusion TP: diffusion from scratch Quiz / Exam |

| | Session | Date | Content |
|--------------|---------|---------|---|
| Foundations | 1 | Jan, 28 | Intro to DL TP: micrograd |
| | 2 | Feb, 4 | Fundamentals I: backprop, loss functions TP: bigram, MLP for next character prediction |
| | 3 | Feb, 11 | Fundamentals II: DL architectures TP: tensor-based models |
| Applications | 4 | Feb, 18 | Attention & Transformers TP: GPT from scratch |
| | 5 | Feb, 25 | DL for Computer vision TP: convnets on CIFAR-10 |
| | 6 | Mar, 3 | VAE and Diffusion TP: diffusion from scratch Quiz / Exam |

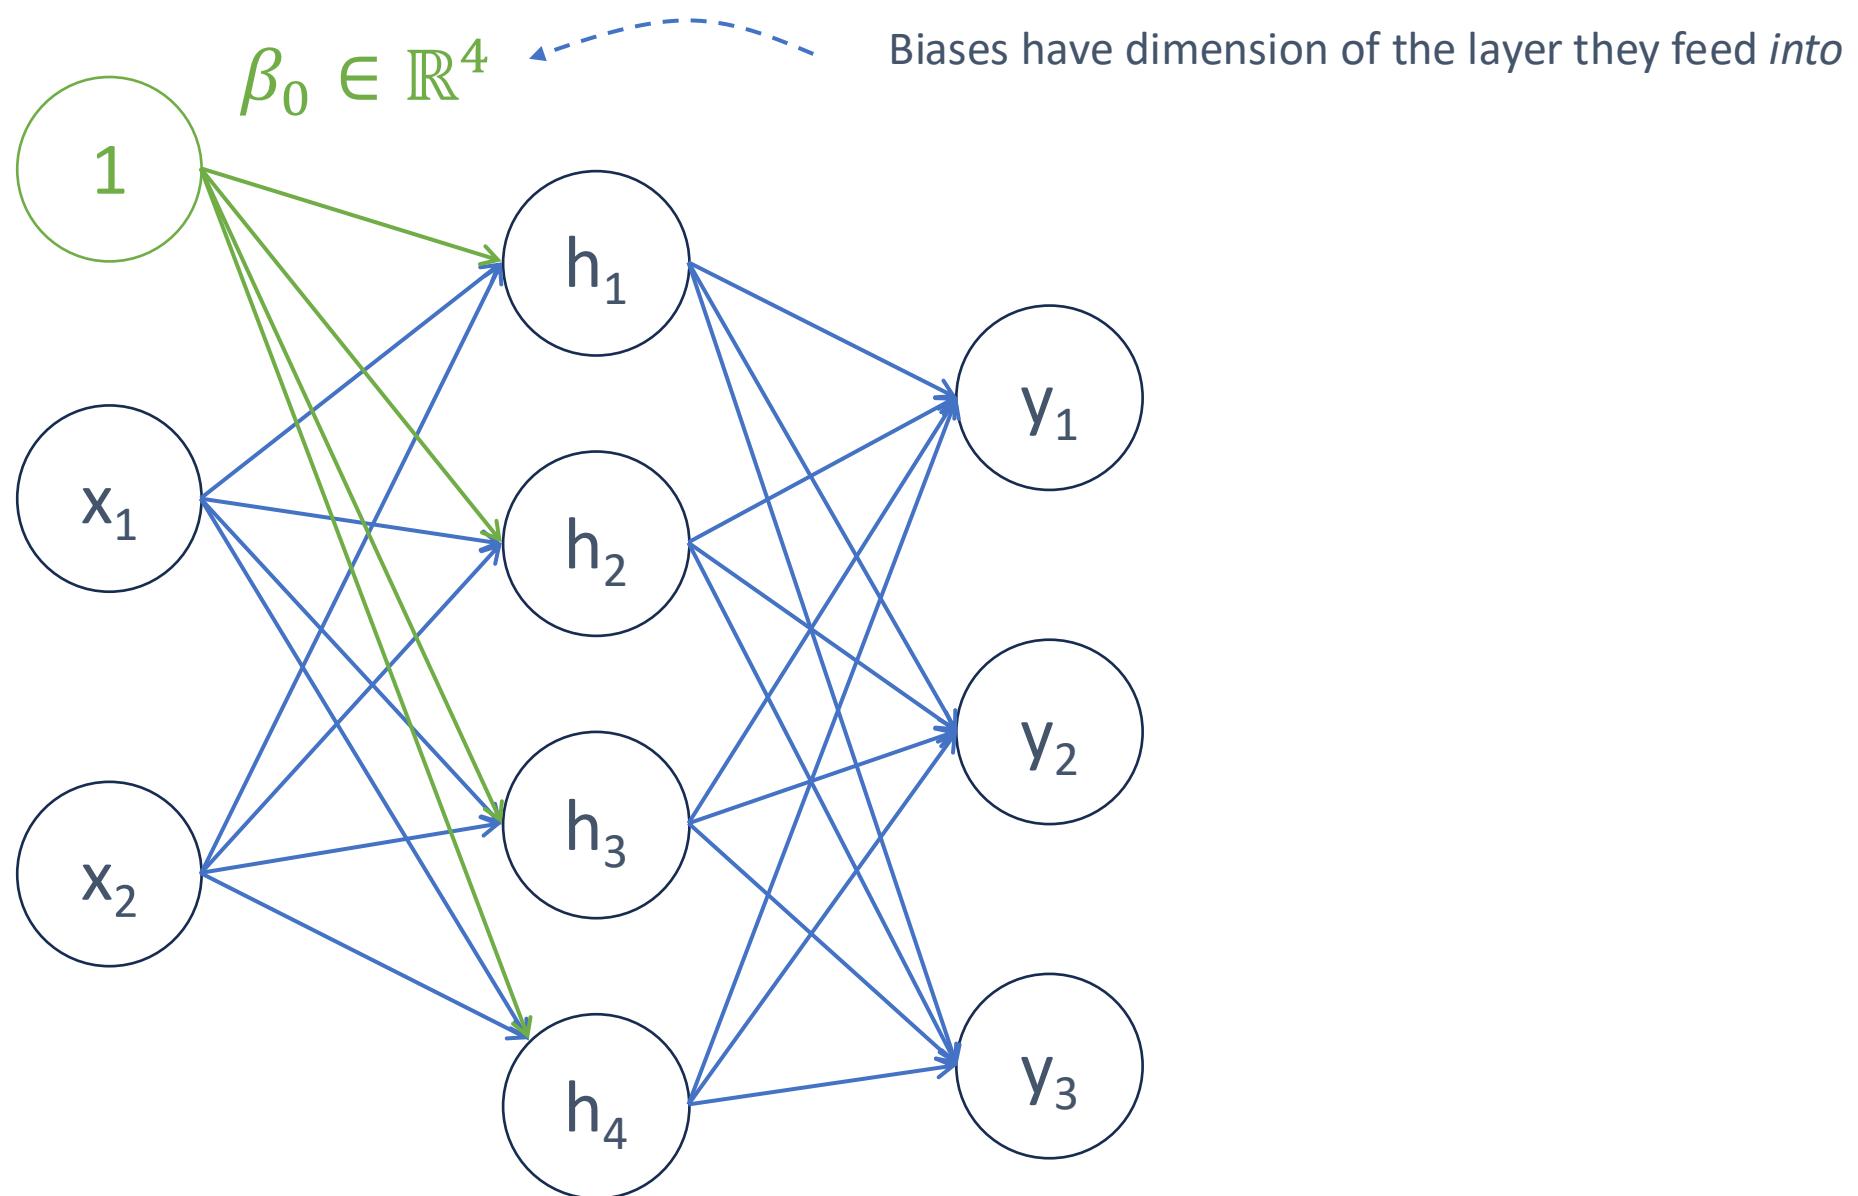


Linear transformation and non-linear activation



$$h_i = a(\beta_i + \sum \theta_{ij} x_j)$$

$$y_i = \gamma_i + \sum \varphi_{ij} h_j$$



Activation functions add non-linearity to the network. Will be discussed in Session 3!

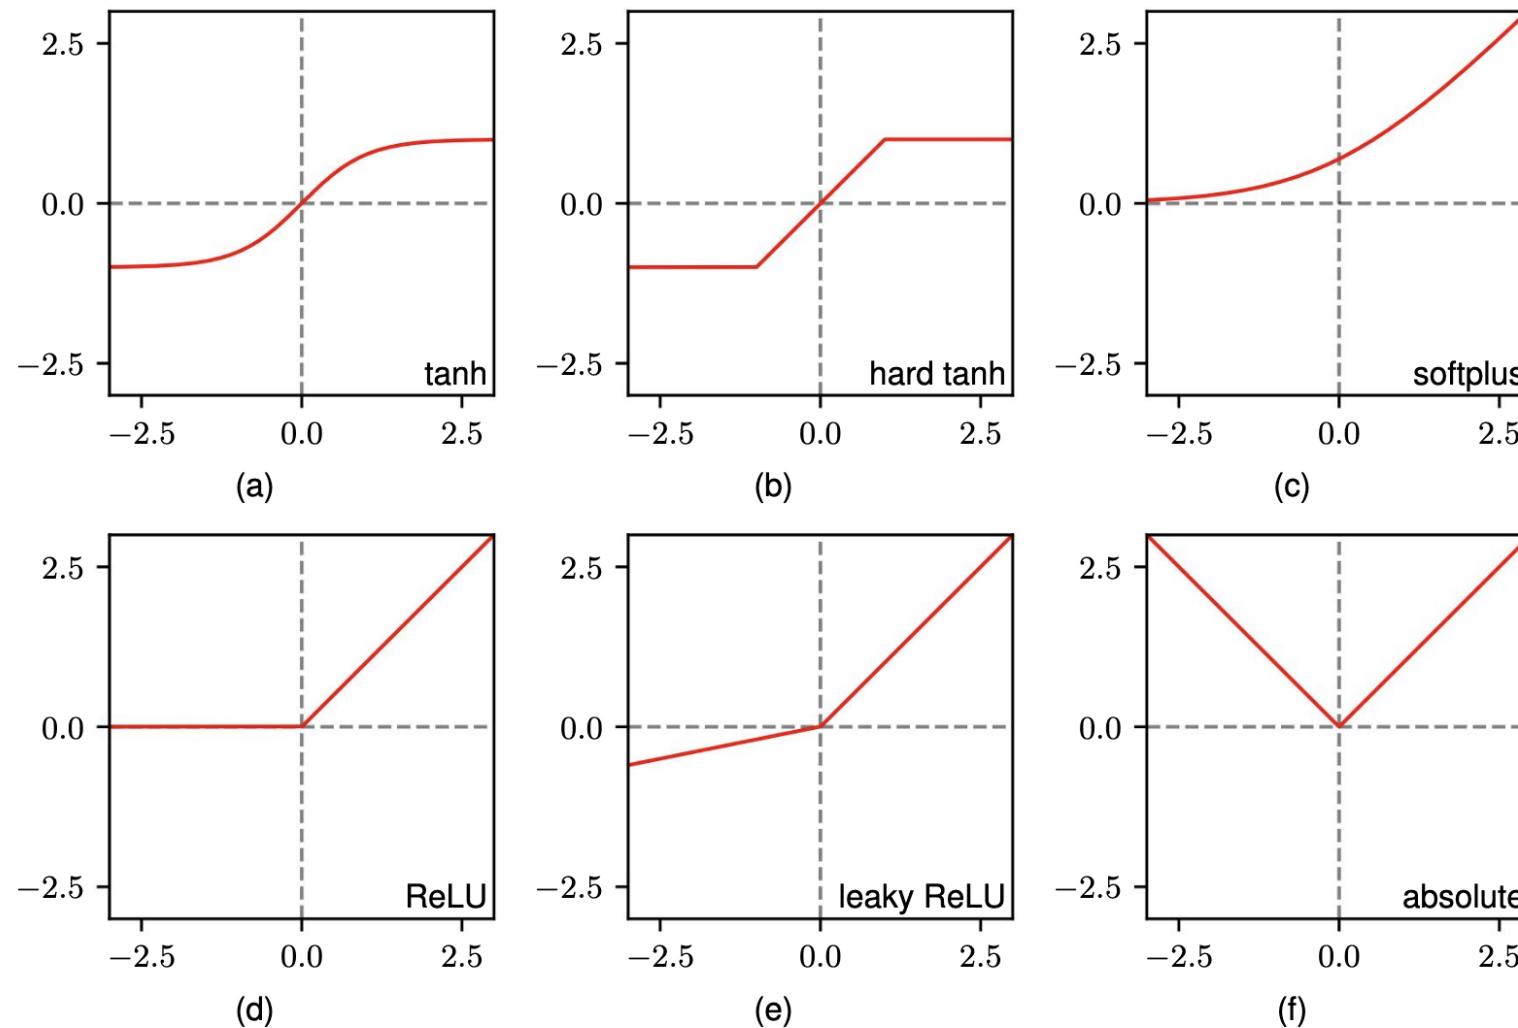
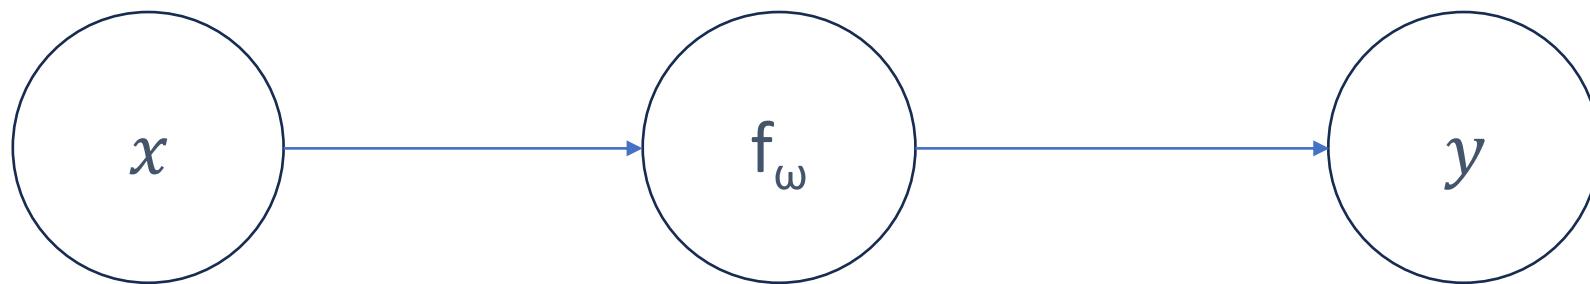


Figure 6.12 A variety of nonlinear activation functions.

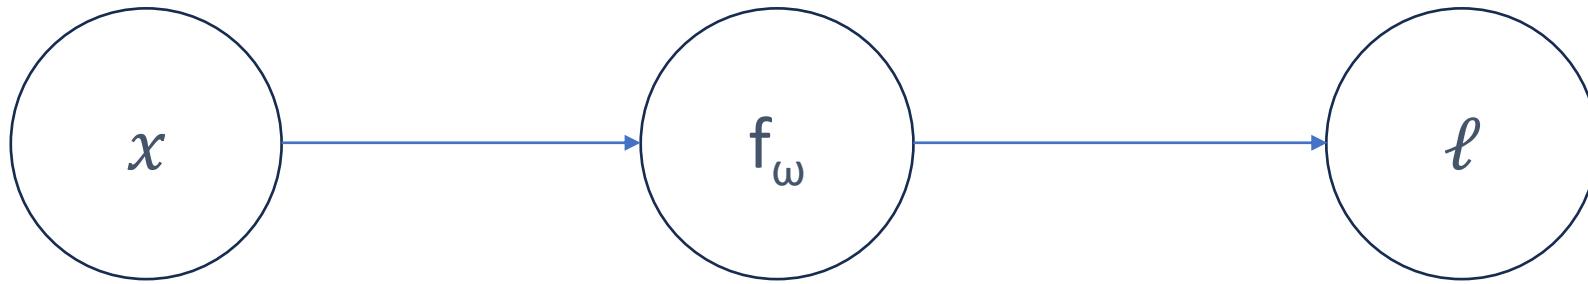
Backprop and gradient descent



$$\ell = (y - \hat{y})^2$$

↑
ground-truth

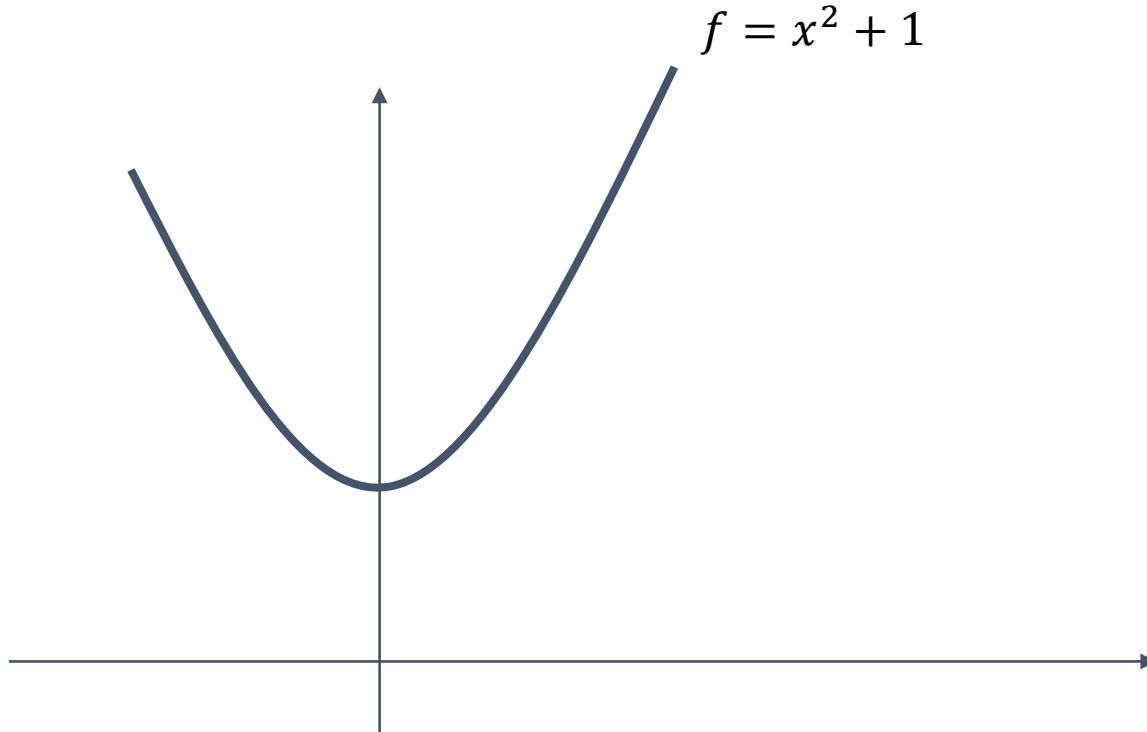
Backprop and gradient descent



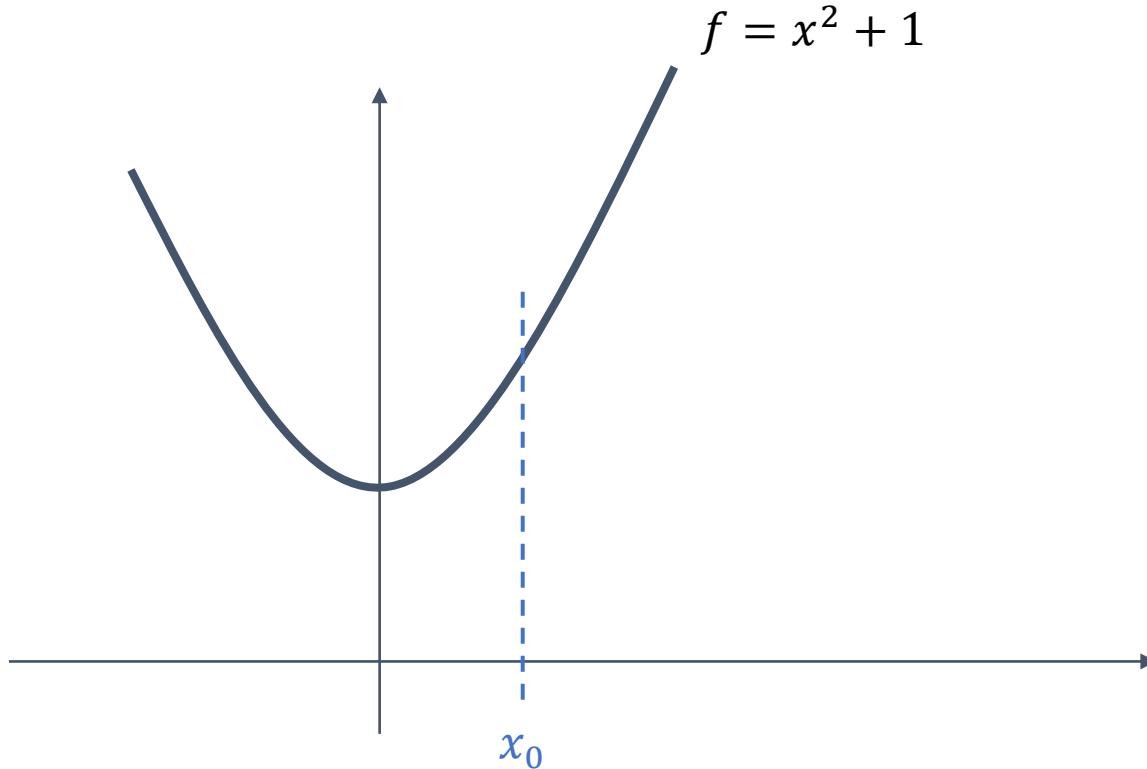
Goal: find ω that minimizes:

$$\ell = f_\omega(x)$$

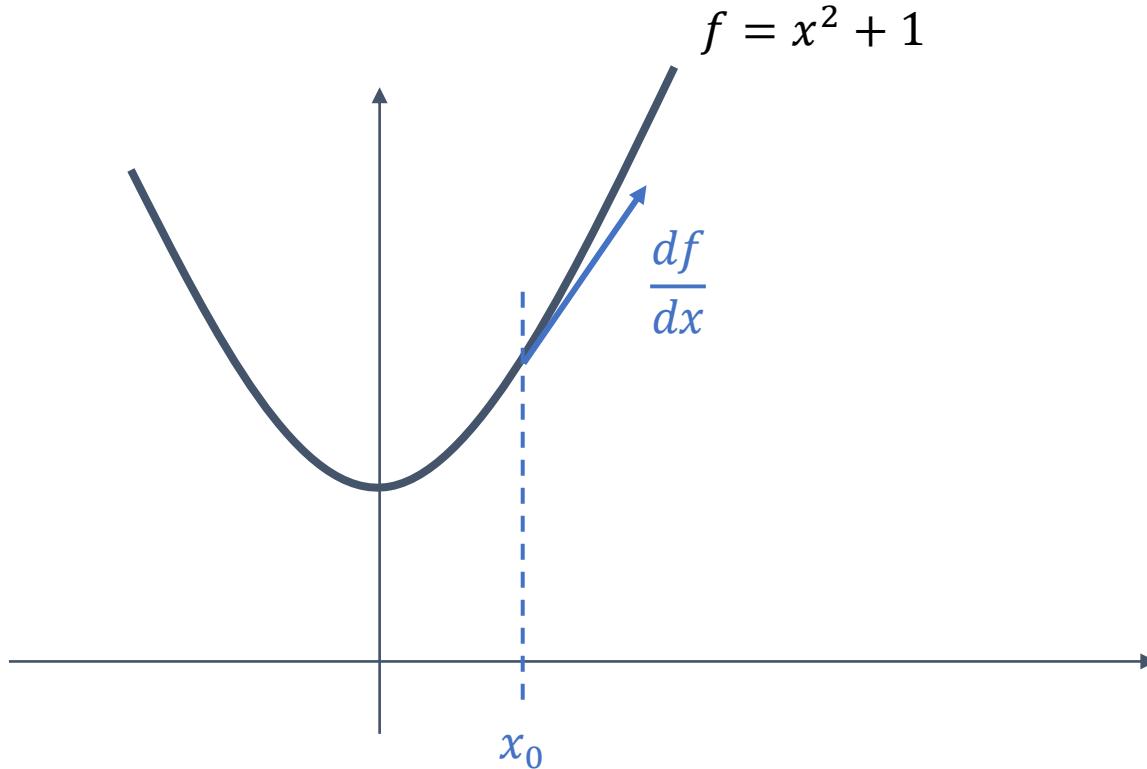
Gradient descent



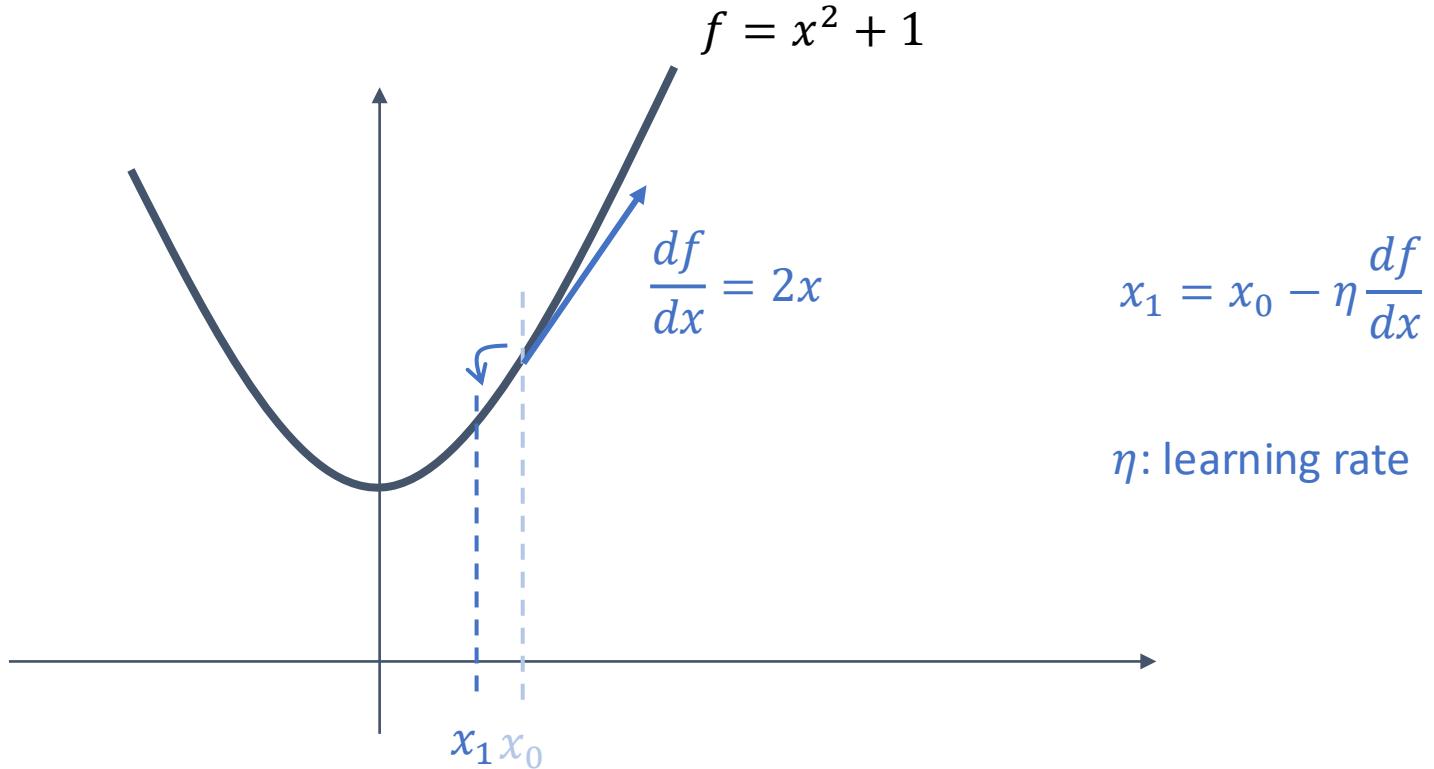
Gradient descent



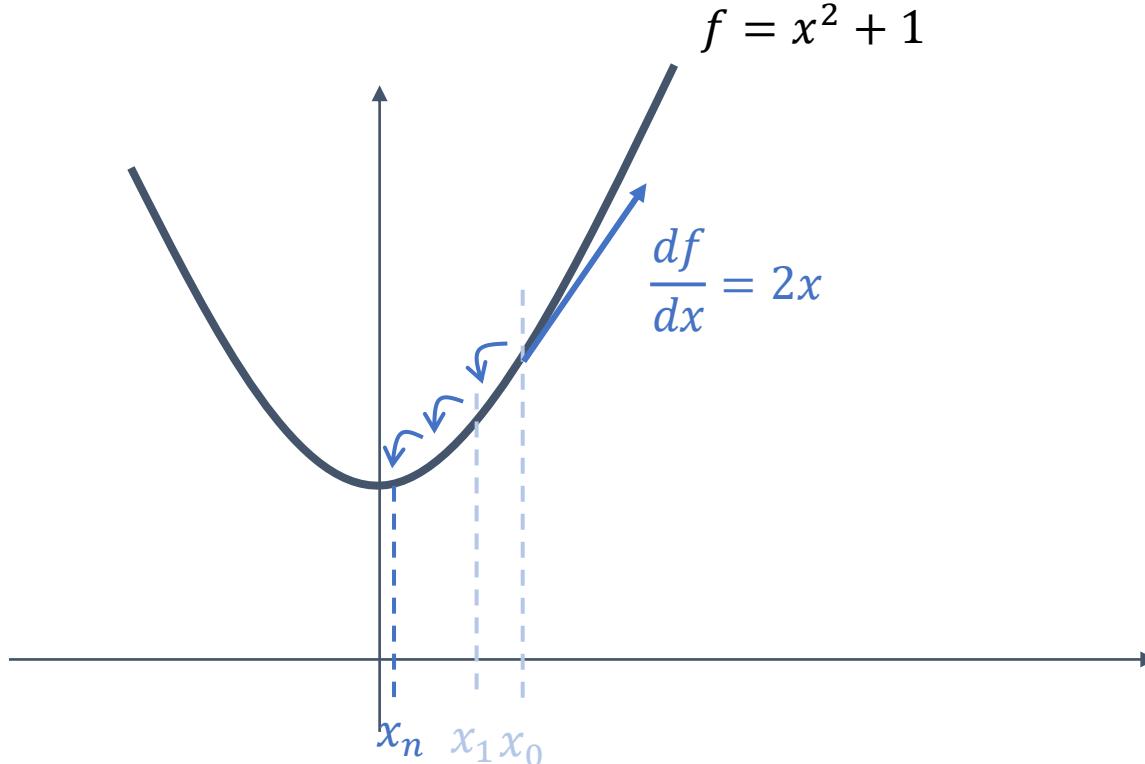
Gradient descent



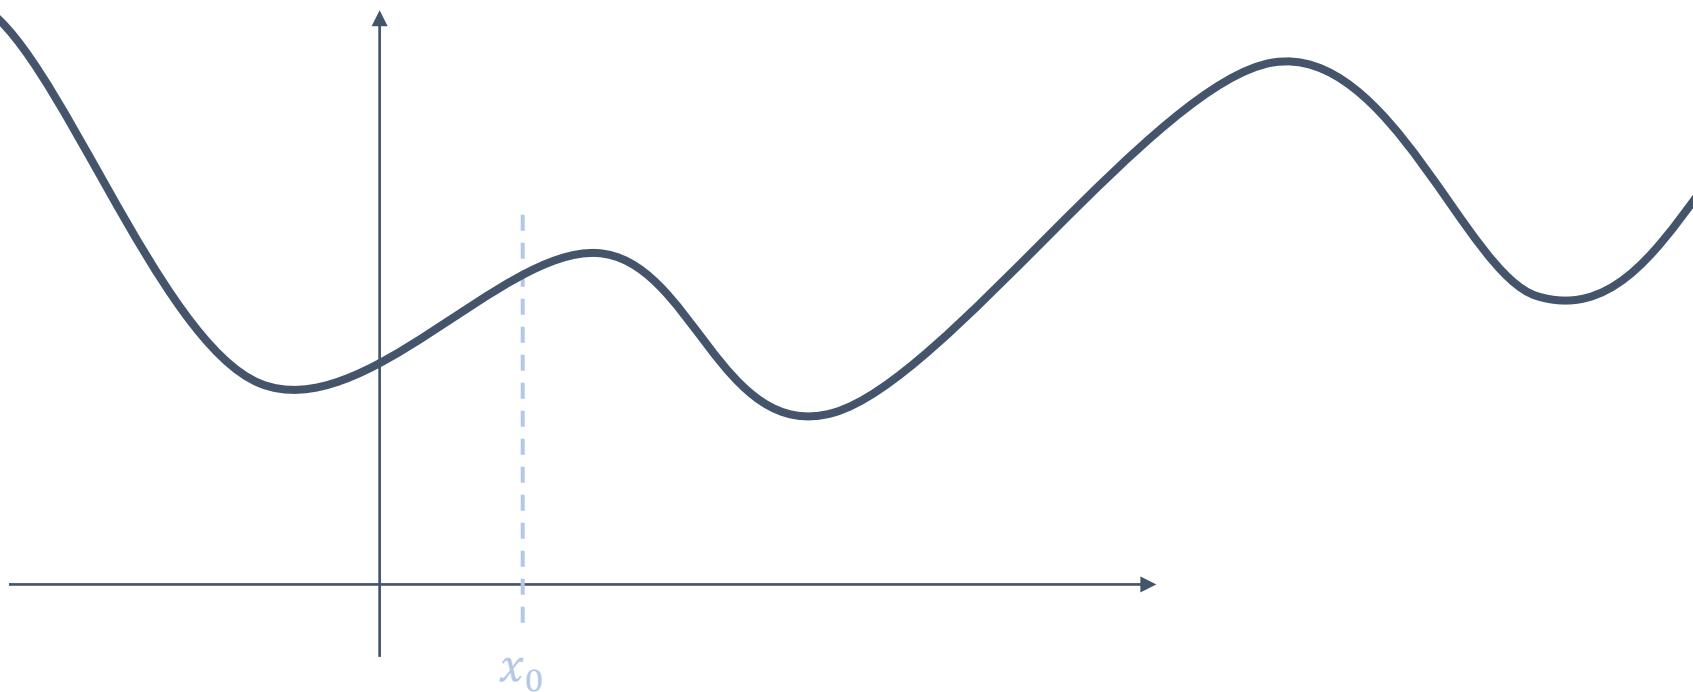
Gradient descent



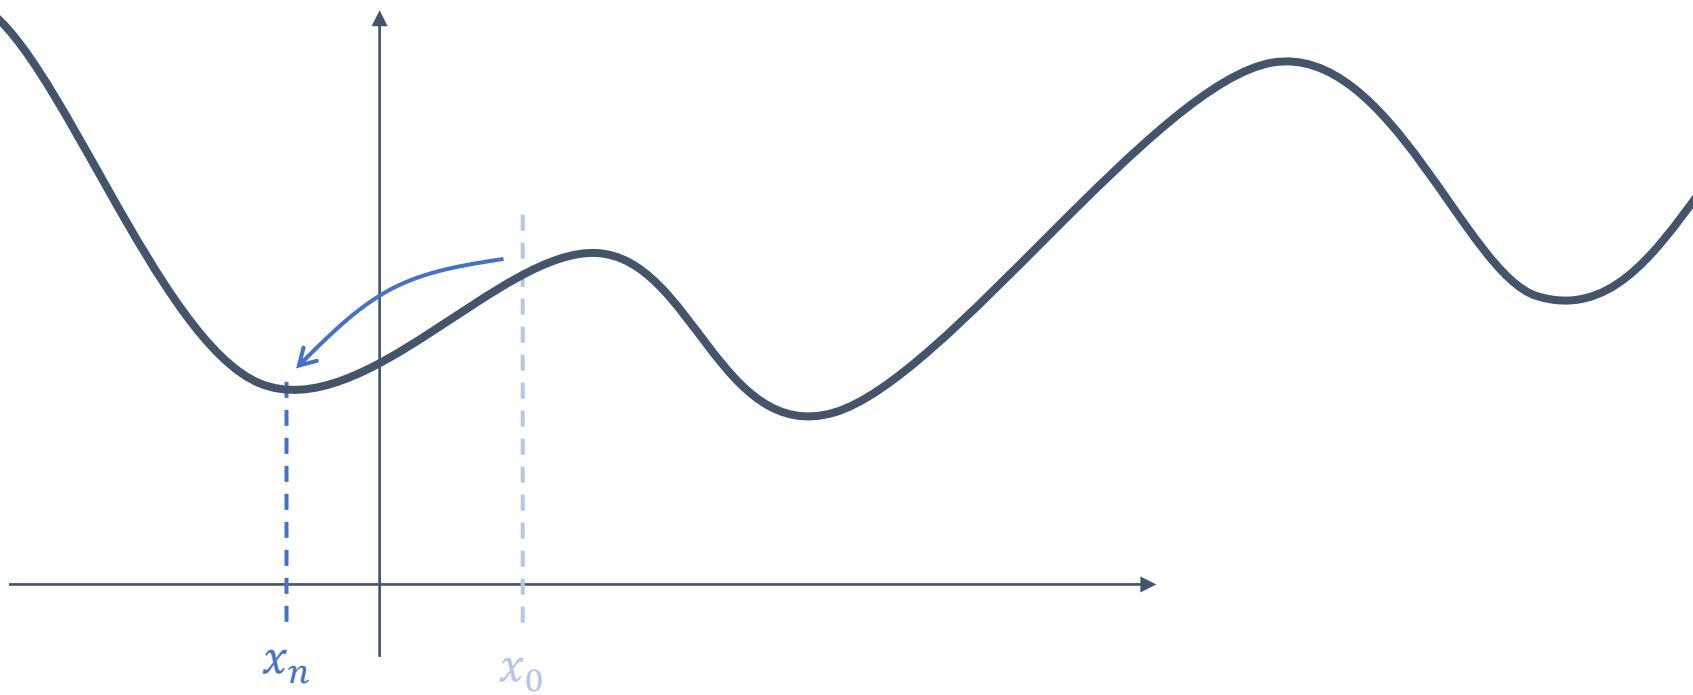
Gradient descent



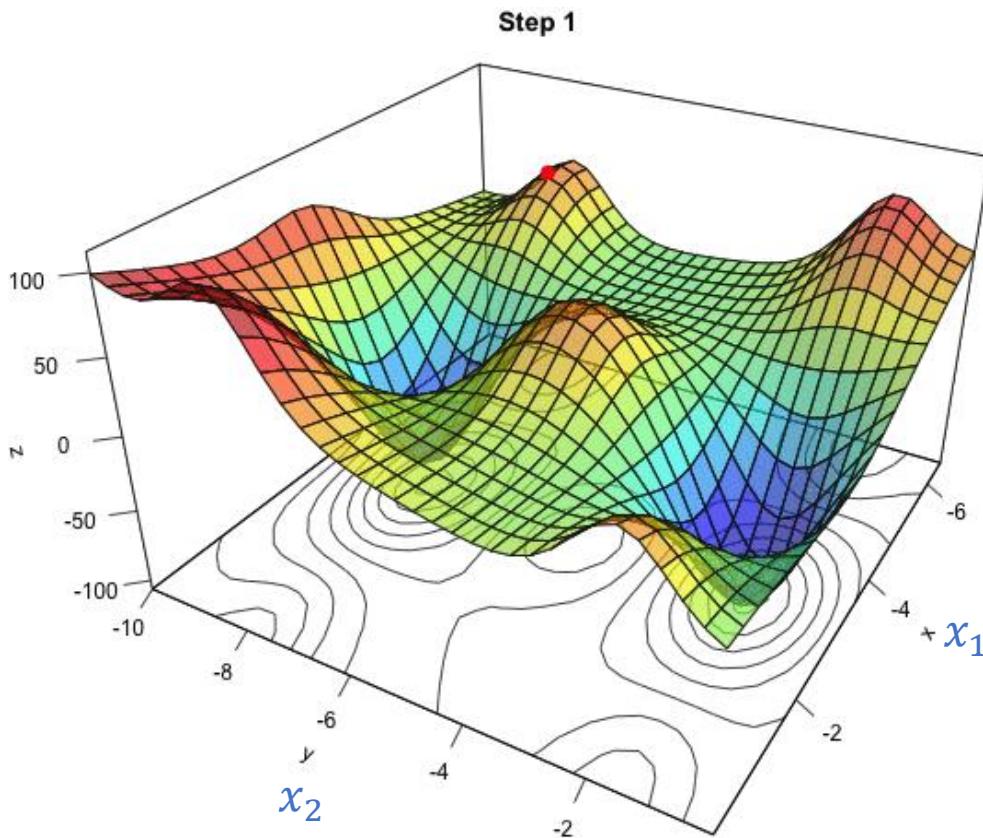
Gradient descent



Gradient descent

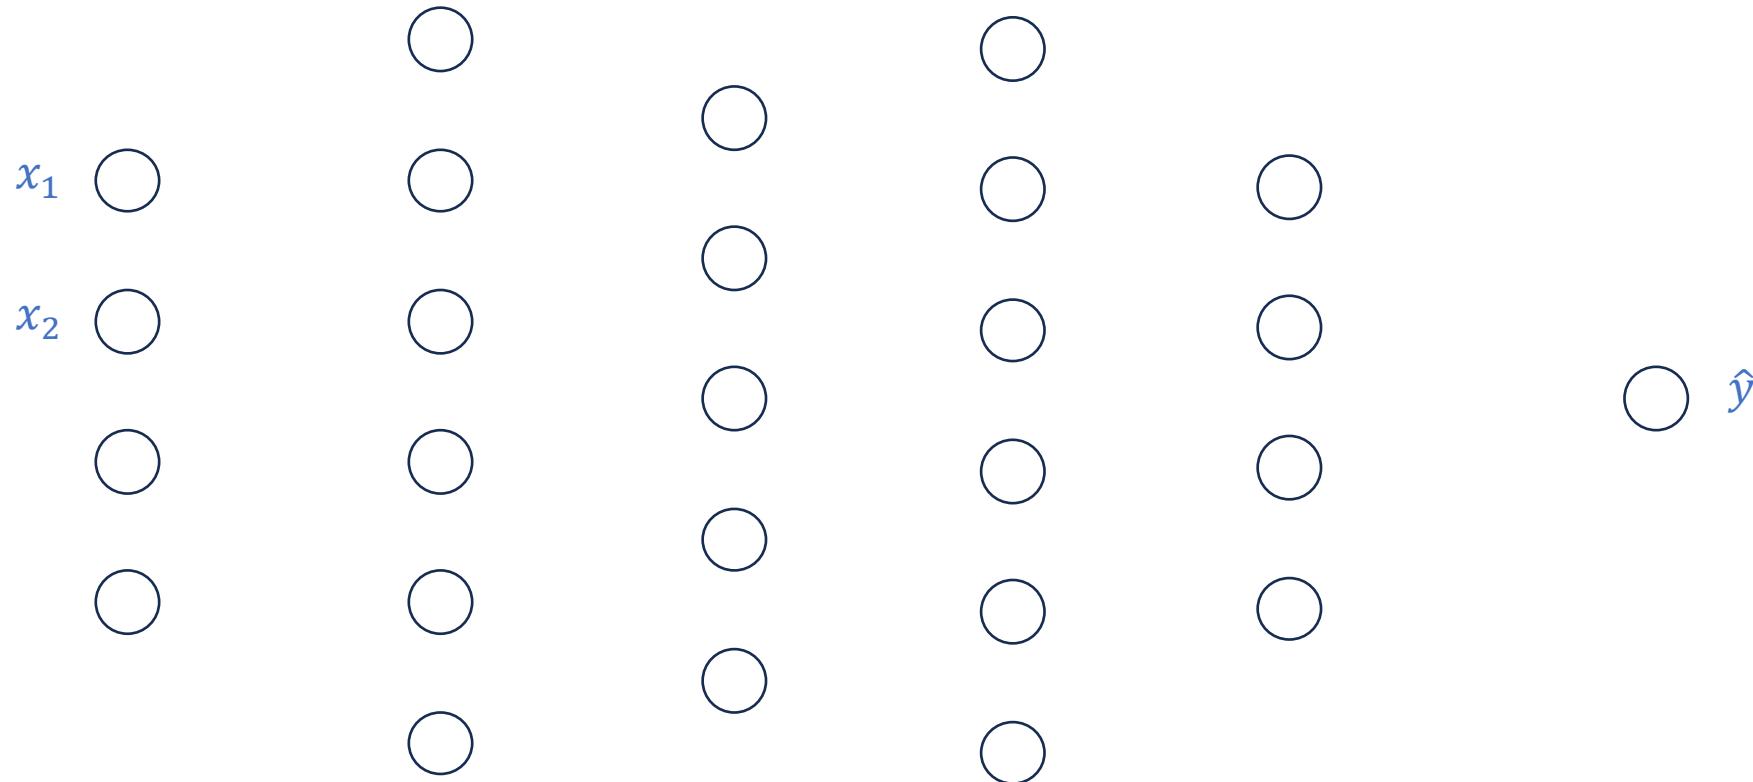


Gradient descent

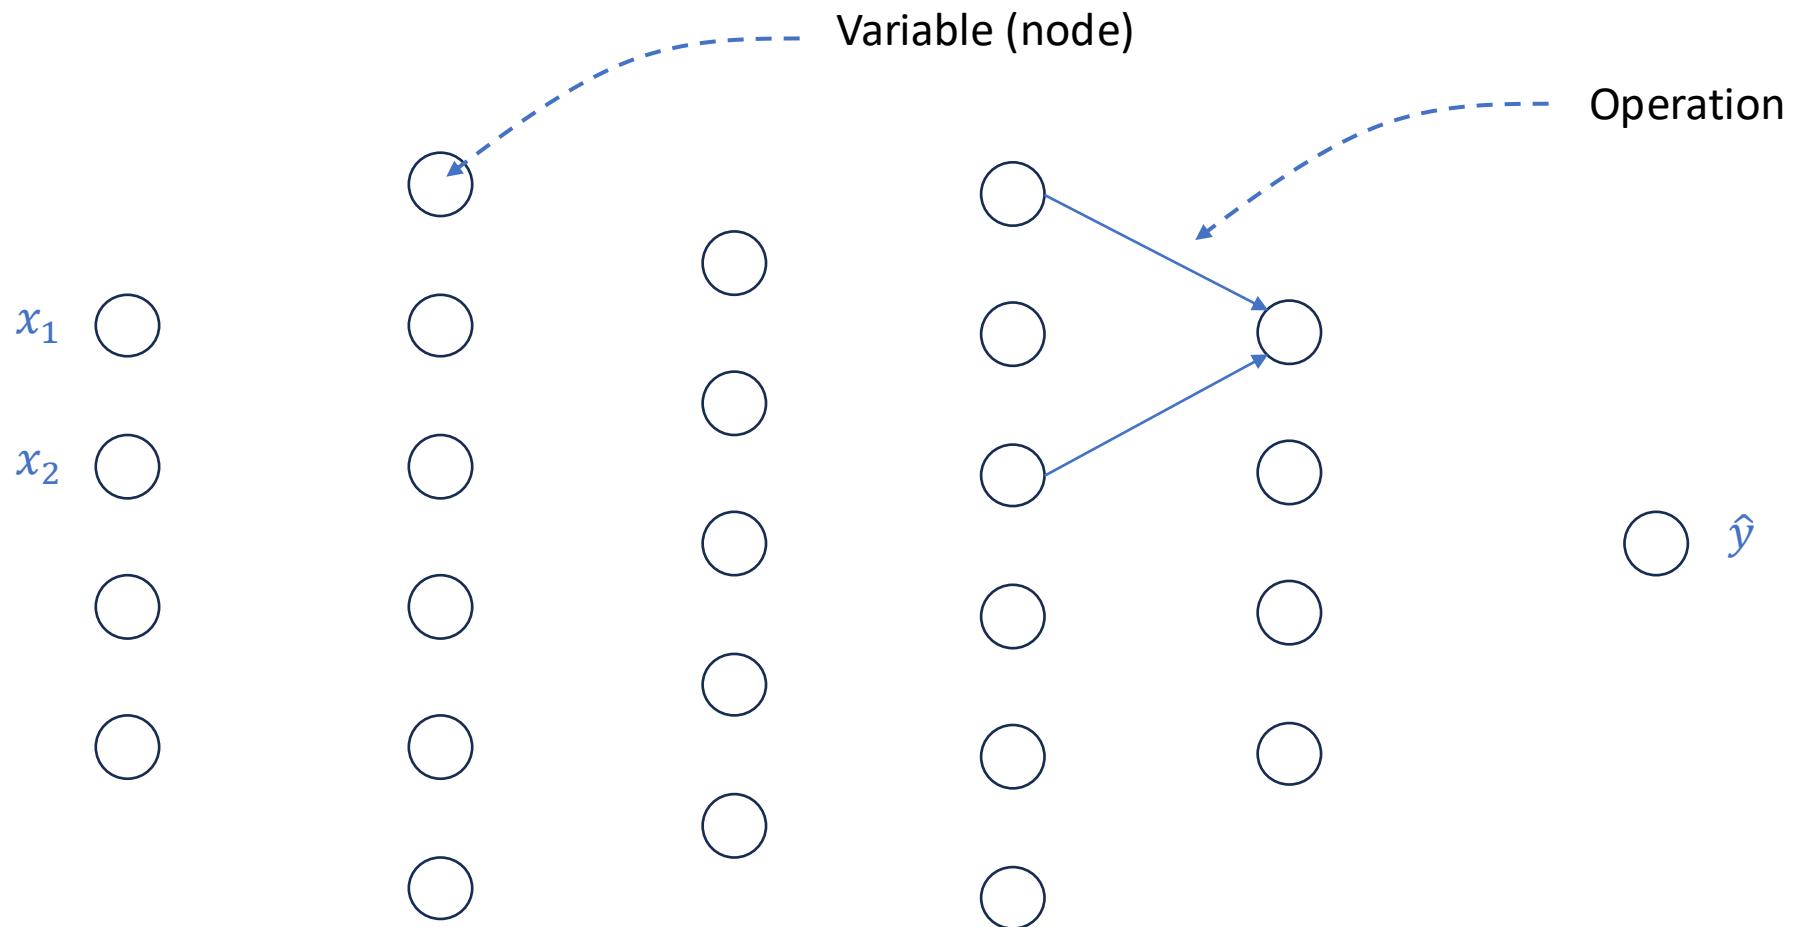


$$\begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \frac{\partial f}{\partial x_2} \end{bmatrix}$$

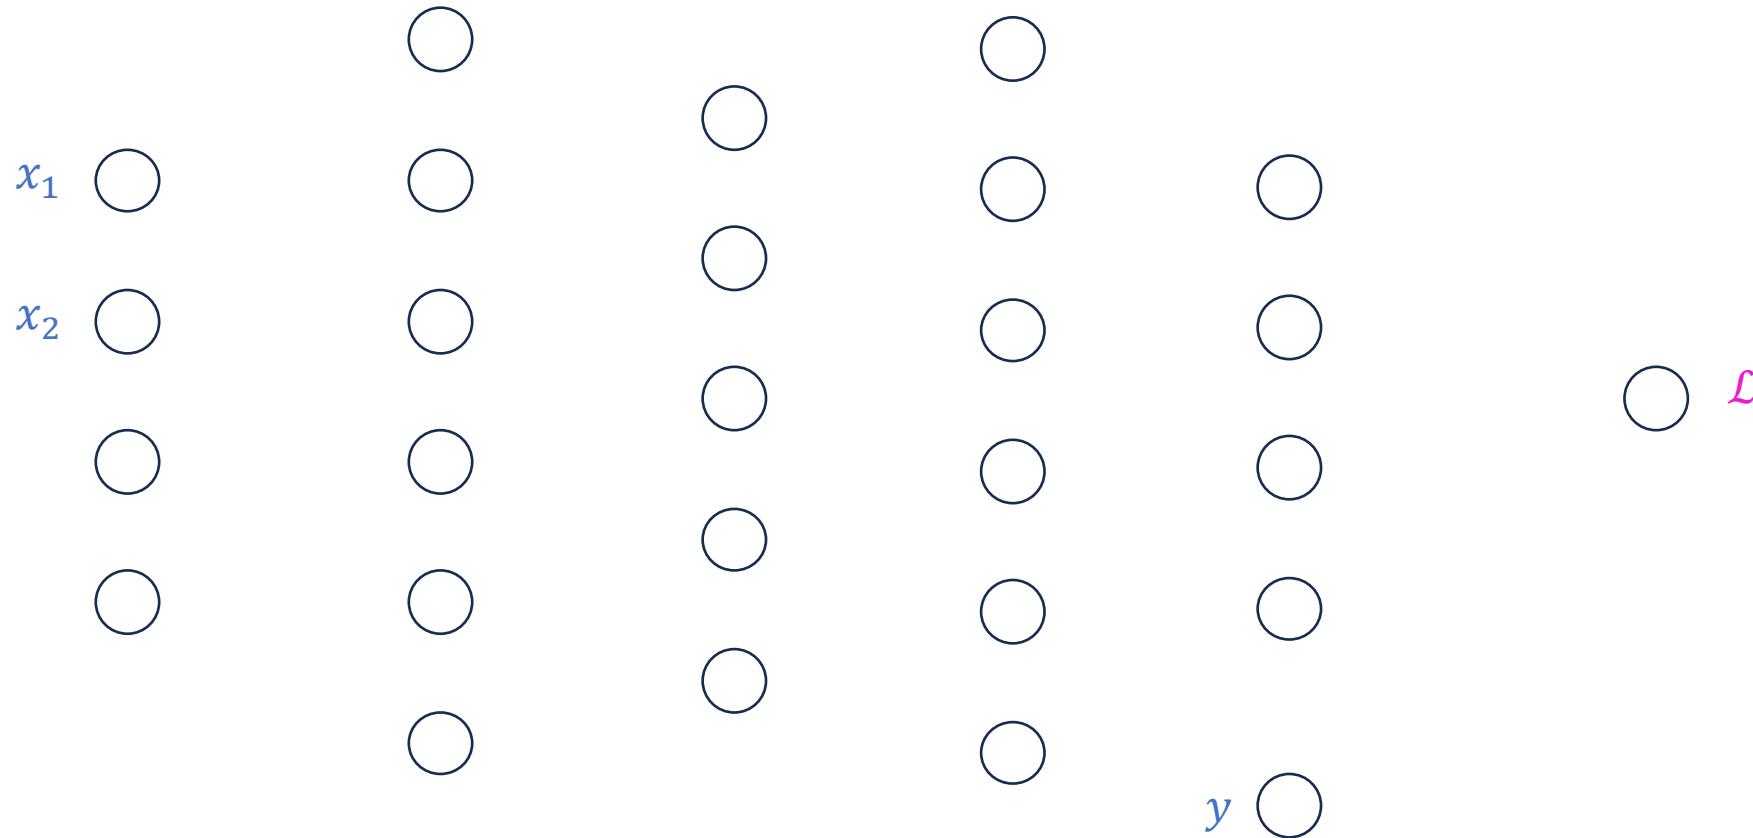
Gradient descent



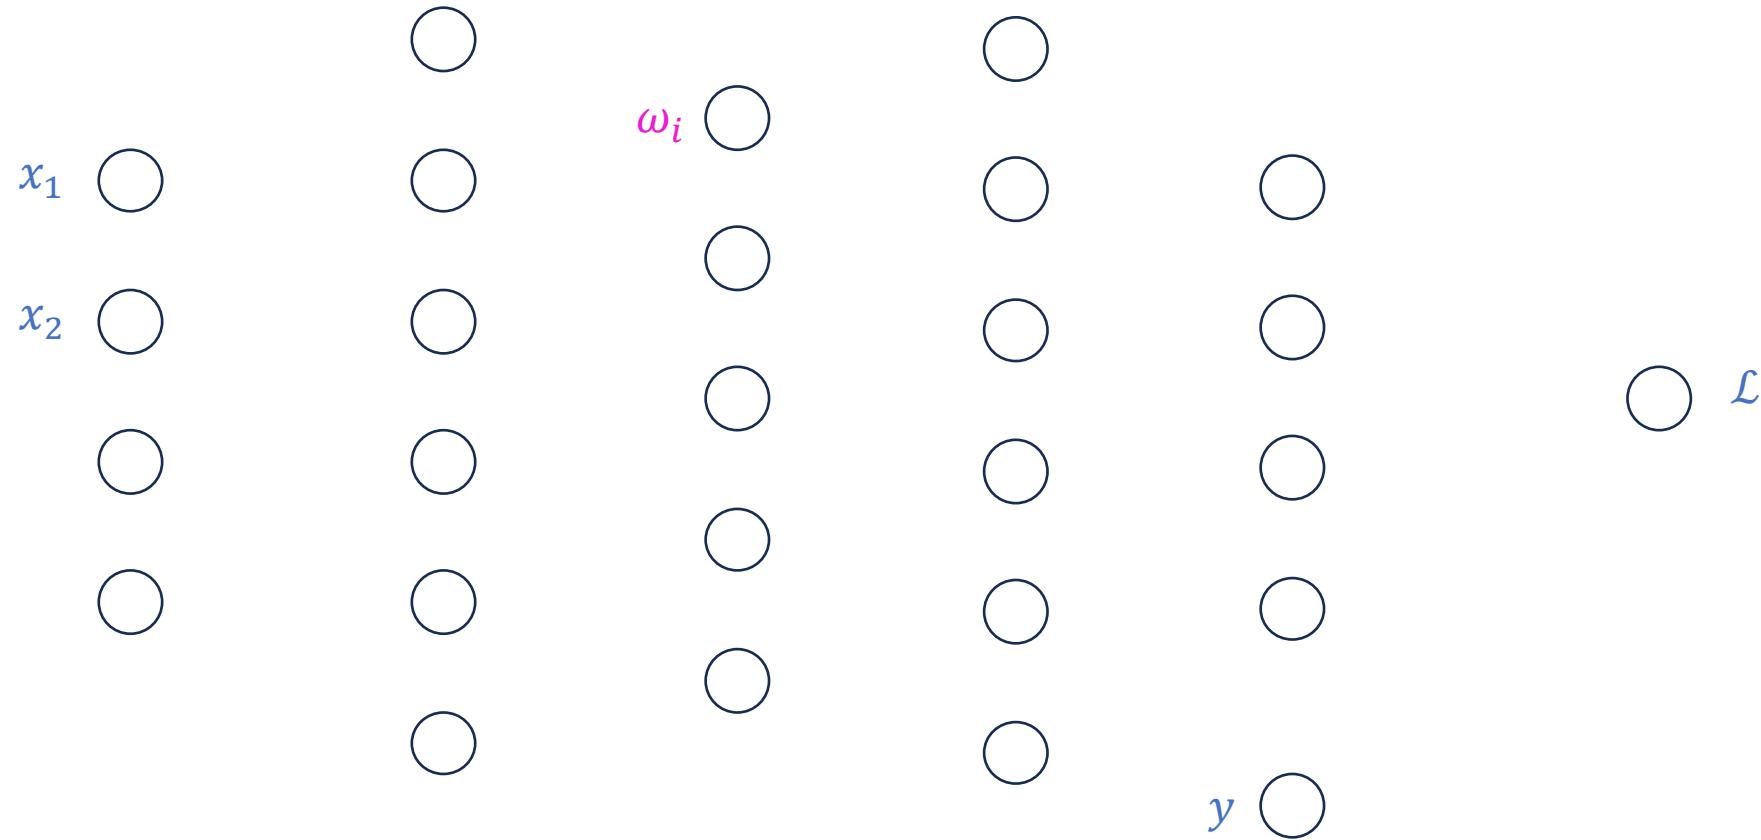
Gradient descent



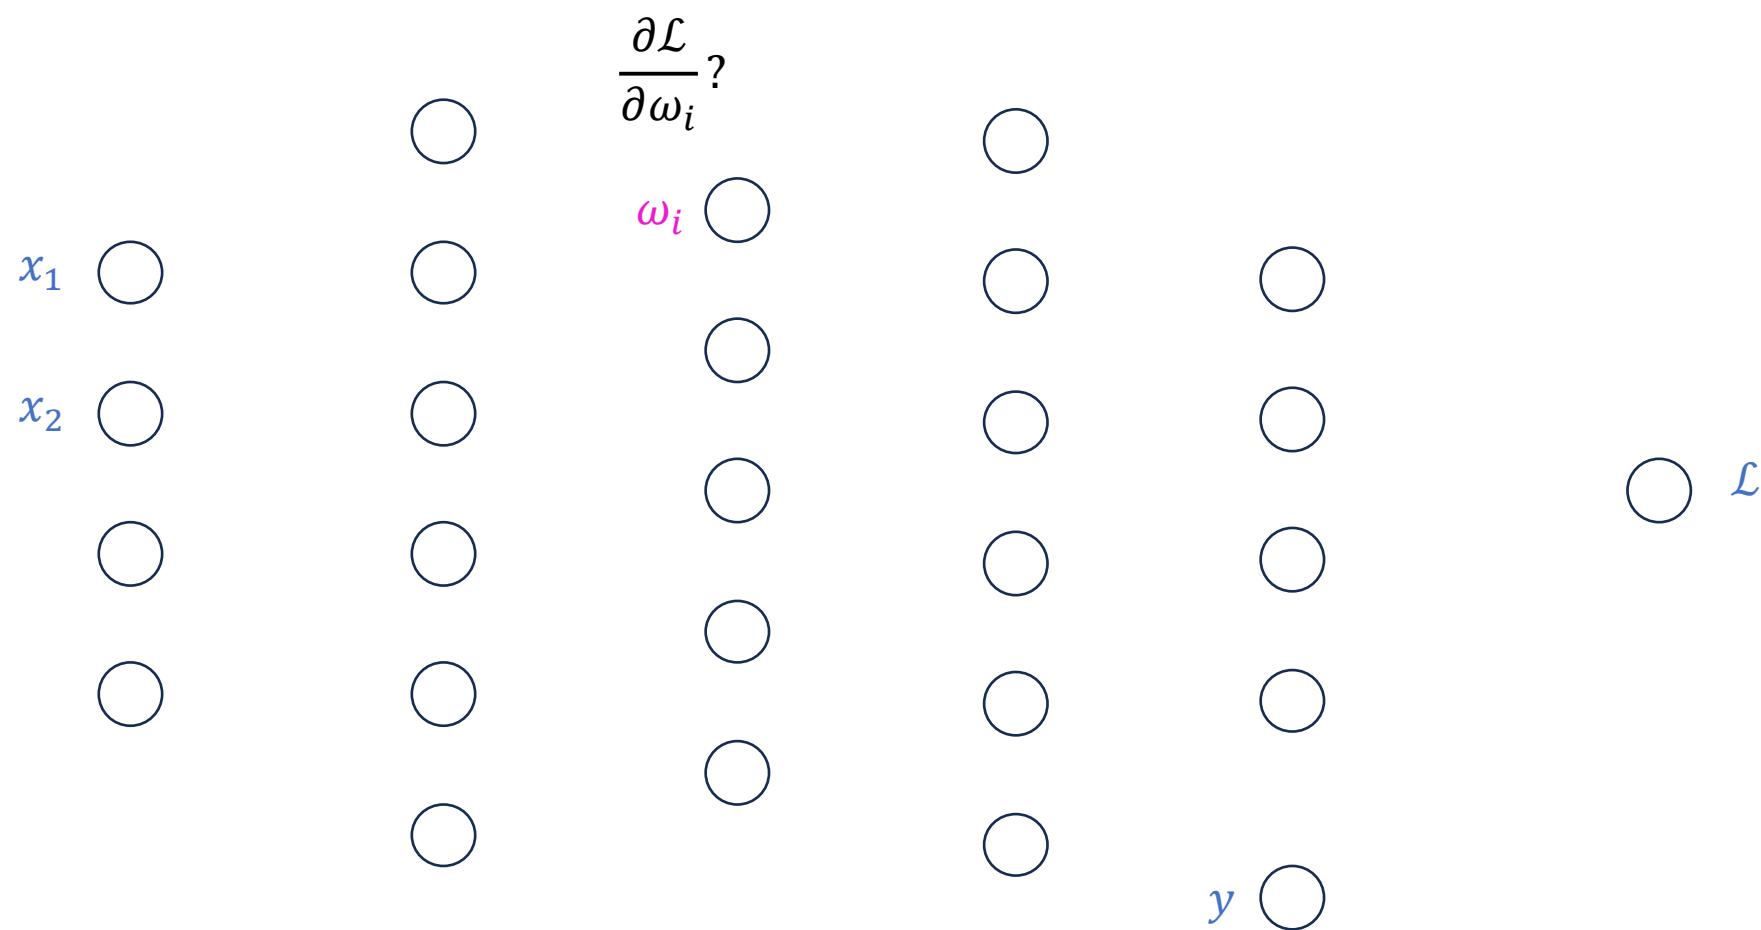
Gradient descent



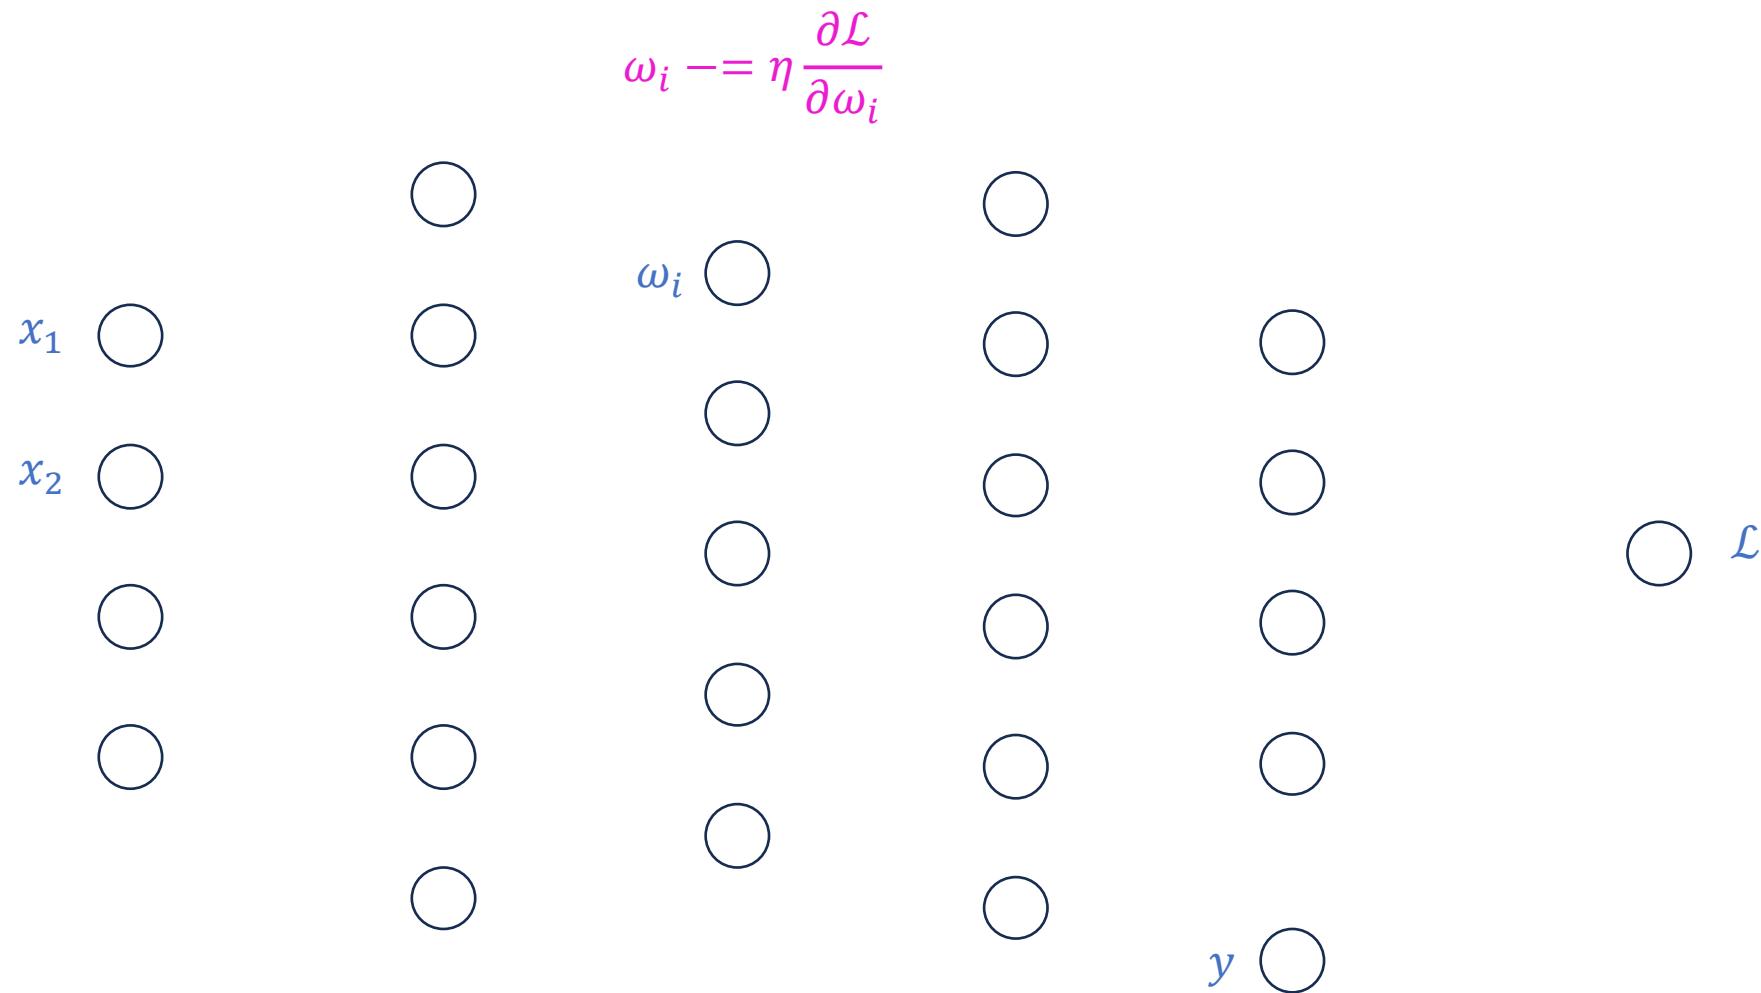
Gradient descent



Gradient descent

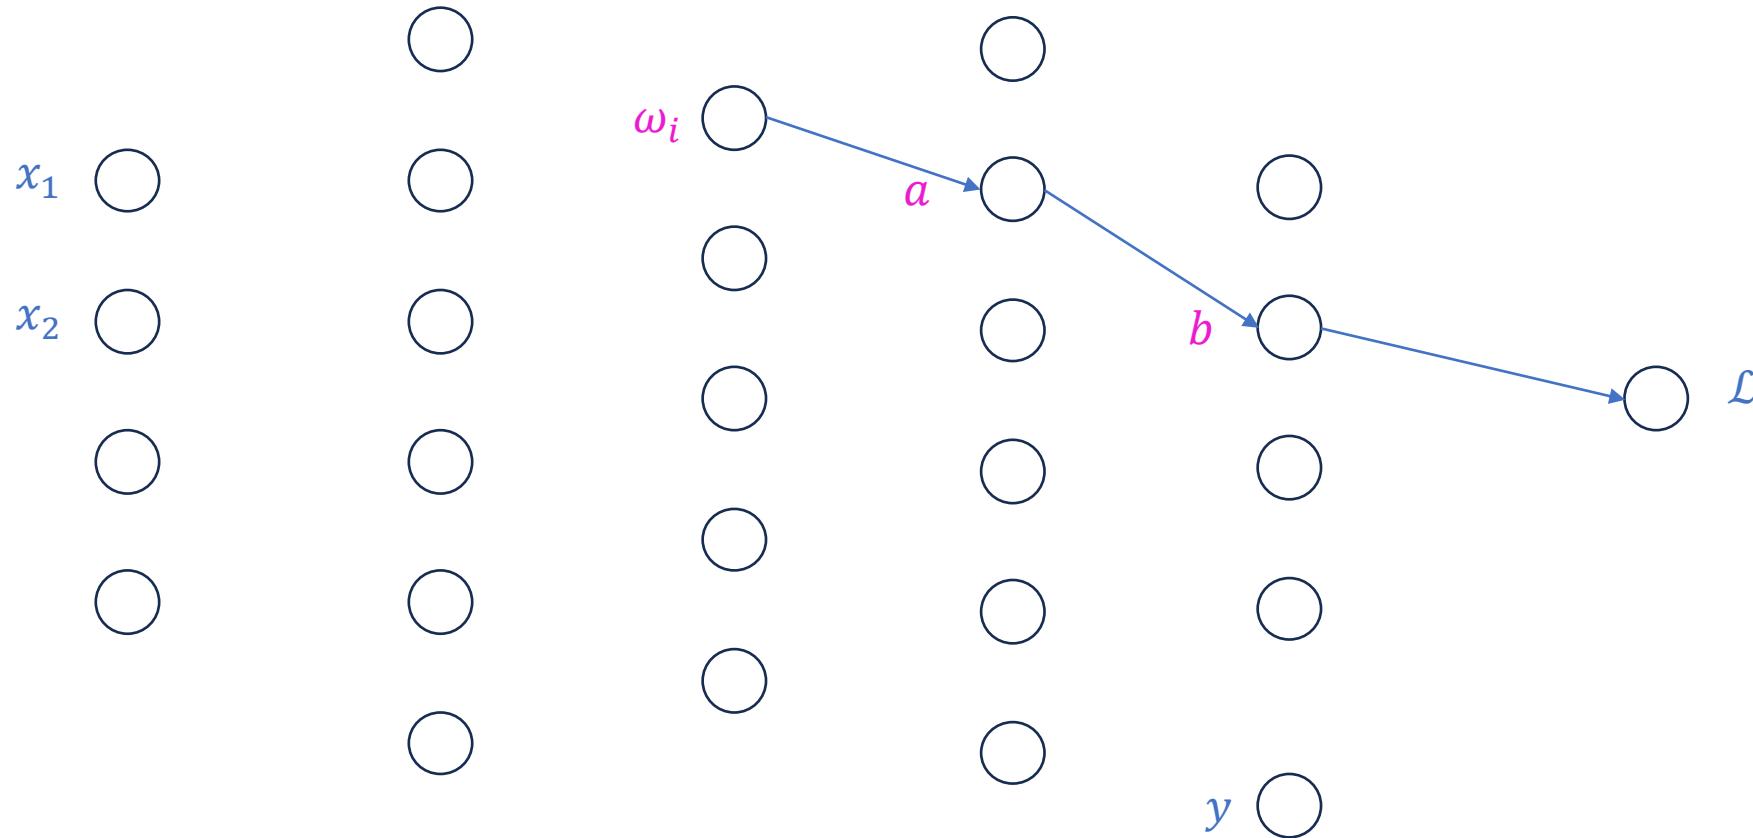


Gradient descent



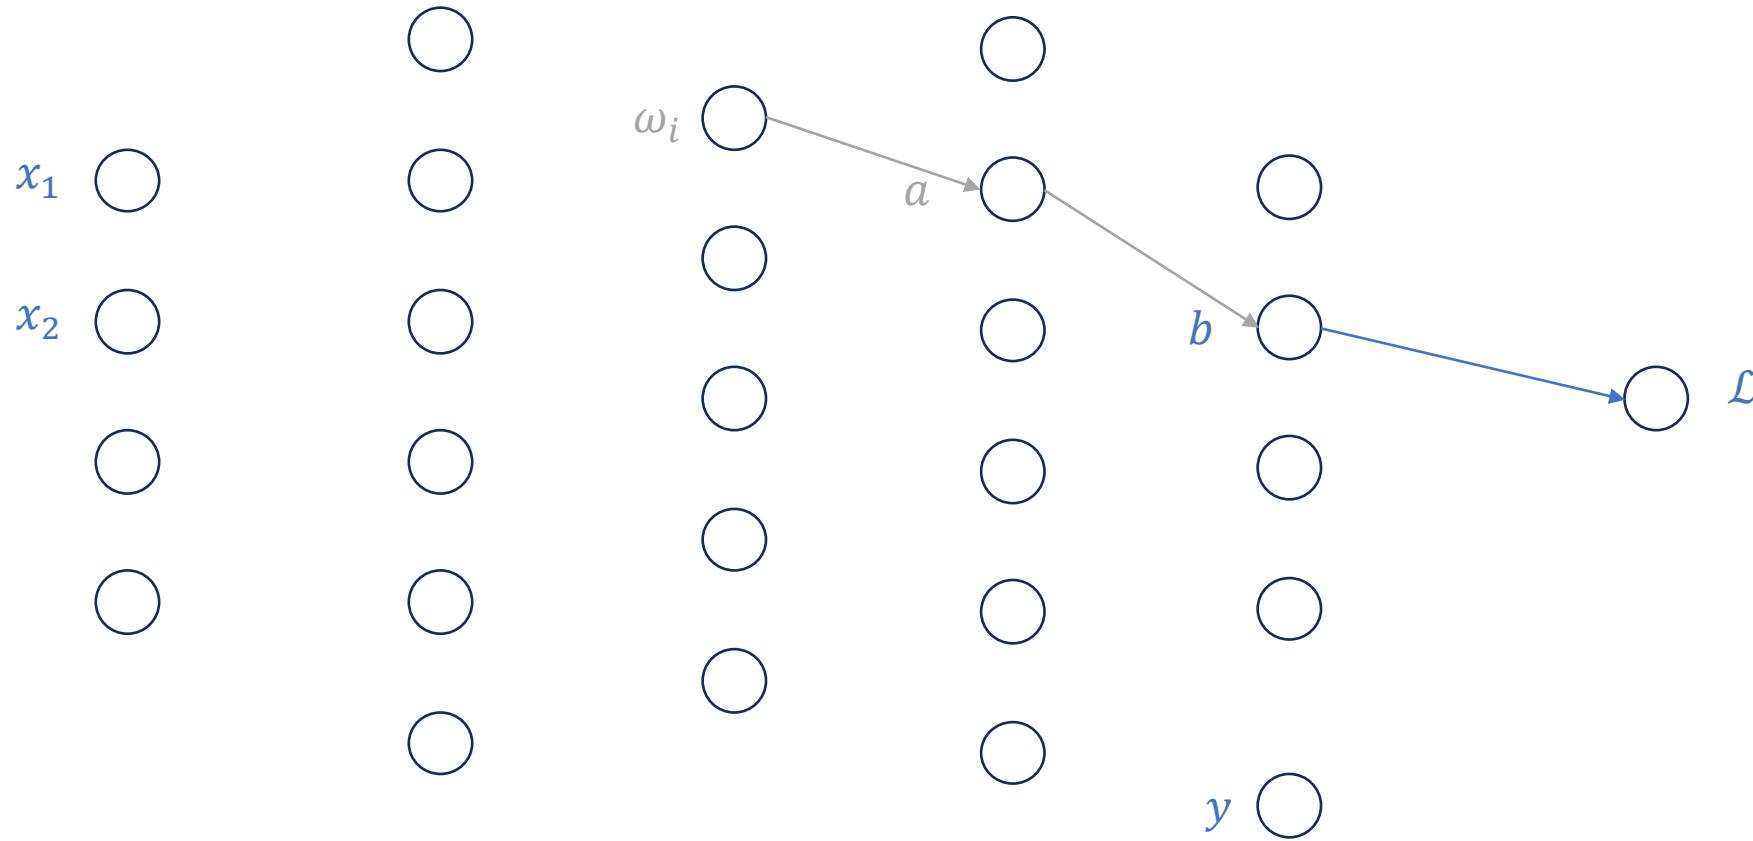
Chain rule

$$\frac{\partial \mathcal{L}}{\partial \omega_i} = \frac{\partial \mathcal{L}}{\partial b} \frac{\partial b}{\partial a} \frac{\partial a}{\partial \omega_i}$$



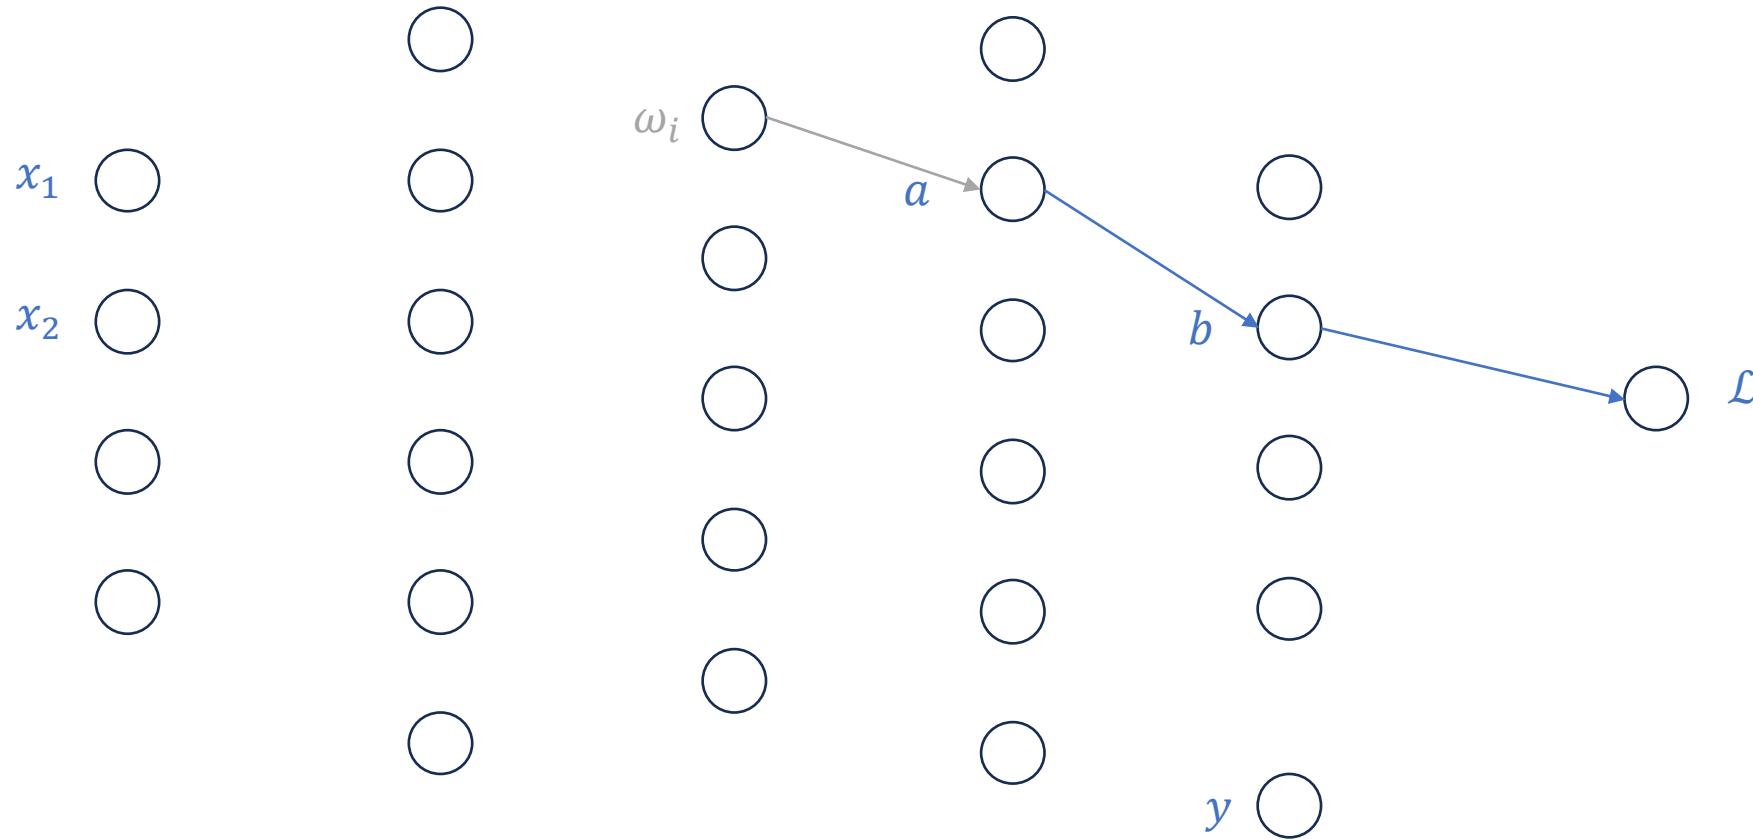
Chain rule

$$\frac{\partial \mathcal{L}}{\partial \omega_i} = \frac{\partial \mathcal{L}}{\partial b} \frac{\partial b}{\partial a} \frac{\partial a}{\partial \omega_i}$$



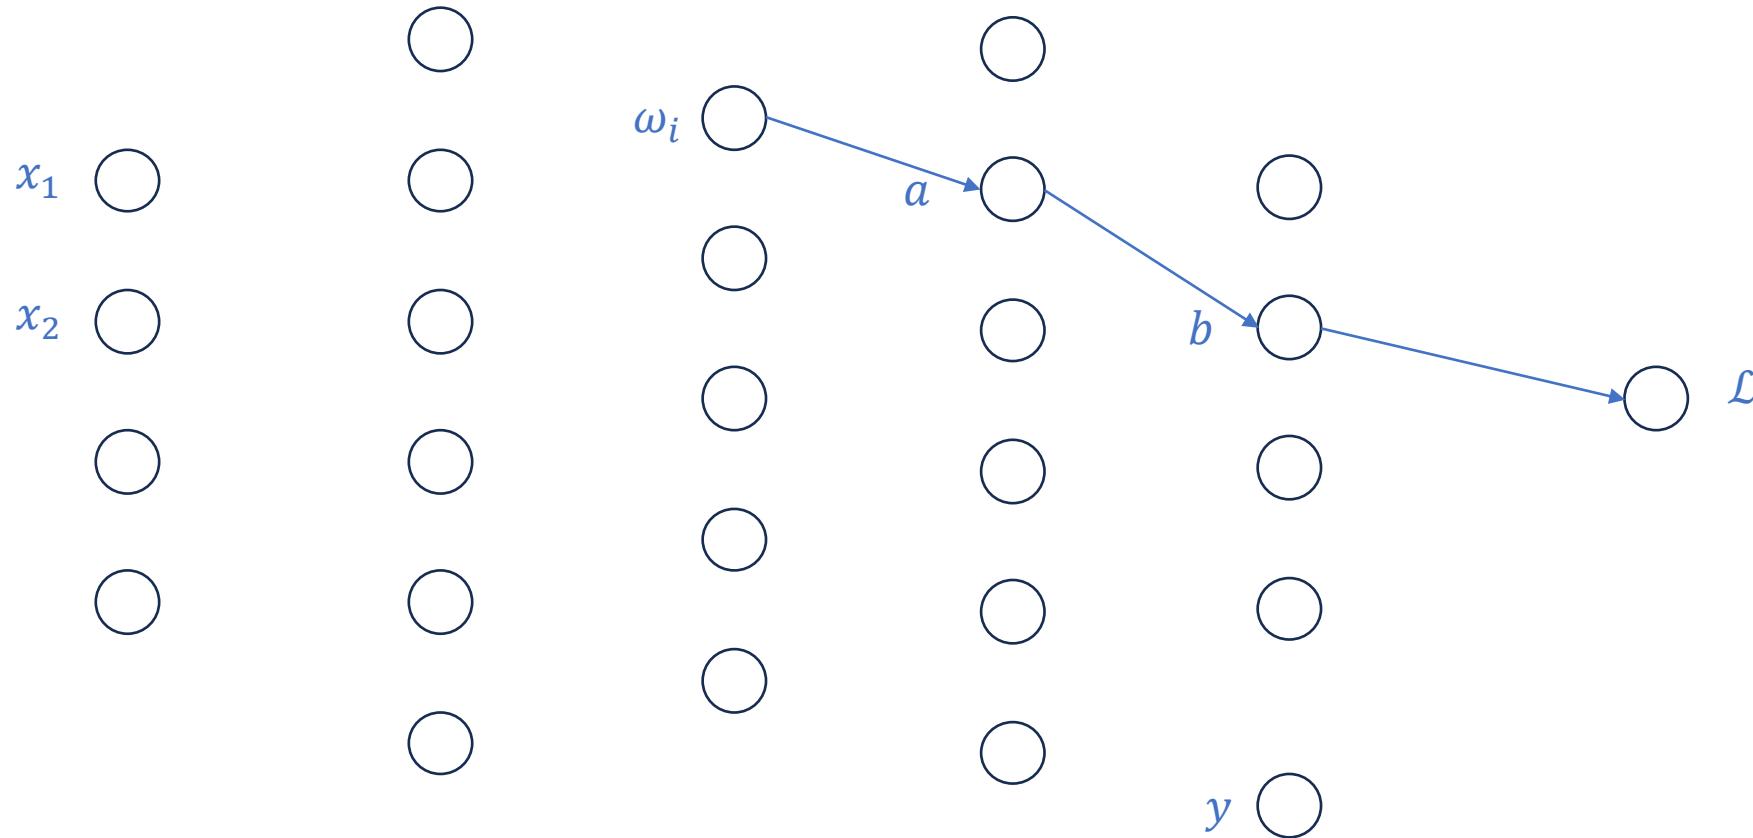
Chain rule

$$\frac{\partial \mathcal{L}}{\partial \omega_i} = \frac{\partial \mathcal{L}}{\partial b} \frac{\partial b}{\partial a} \frac{\partial a}{\partial \omega_i}$$

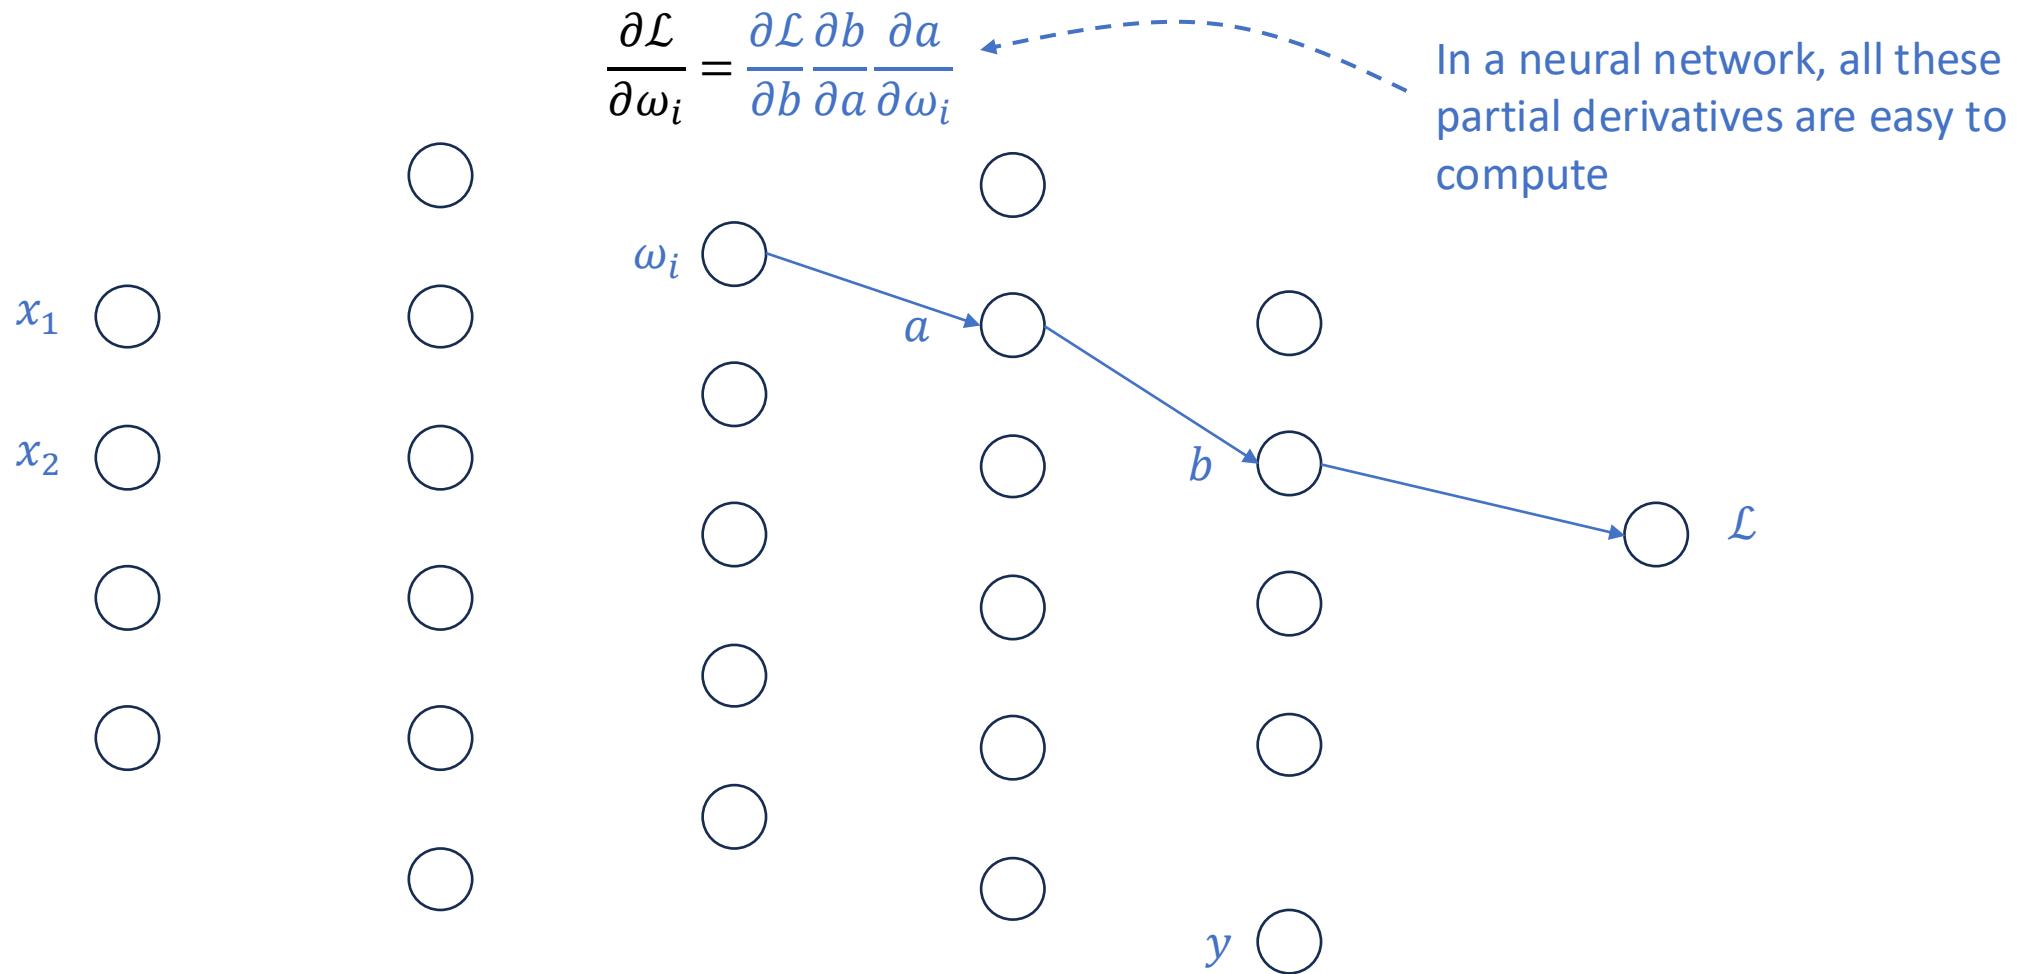


Chain rule

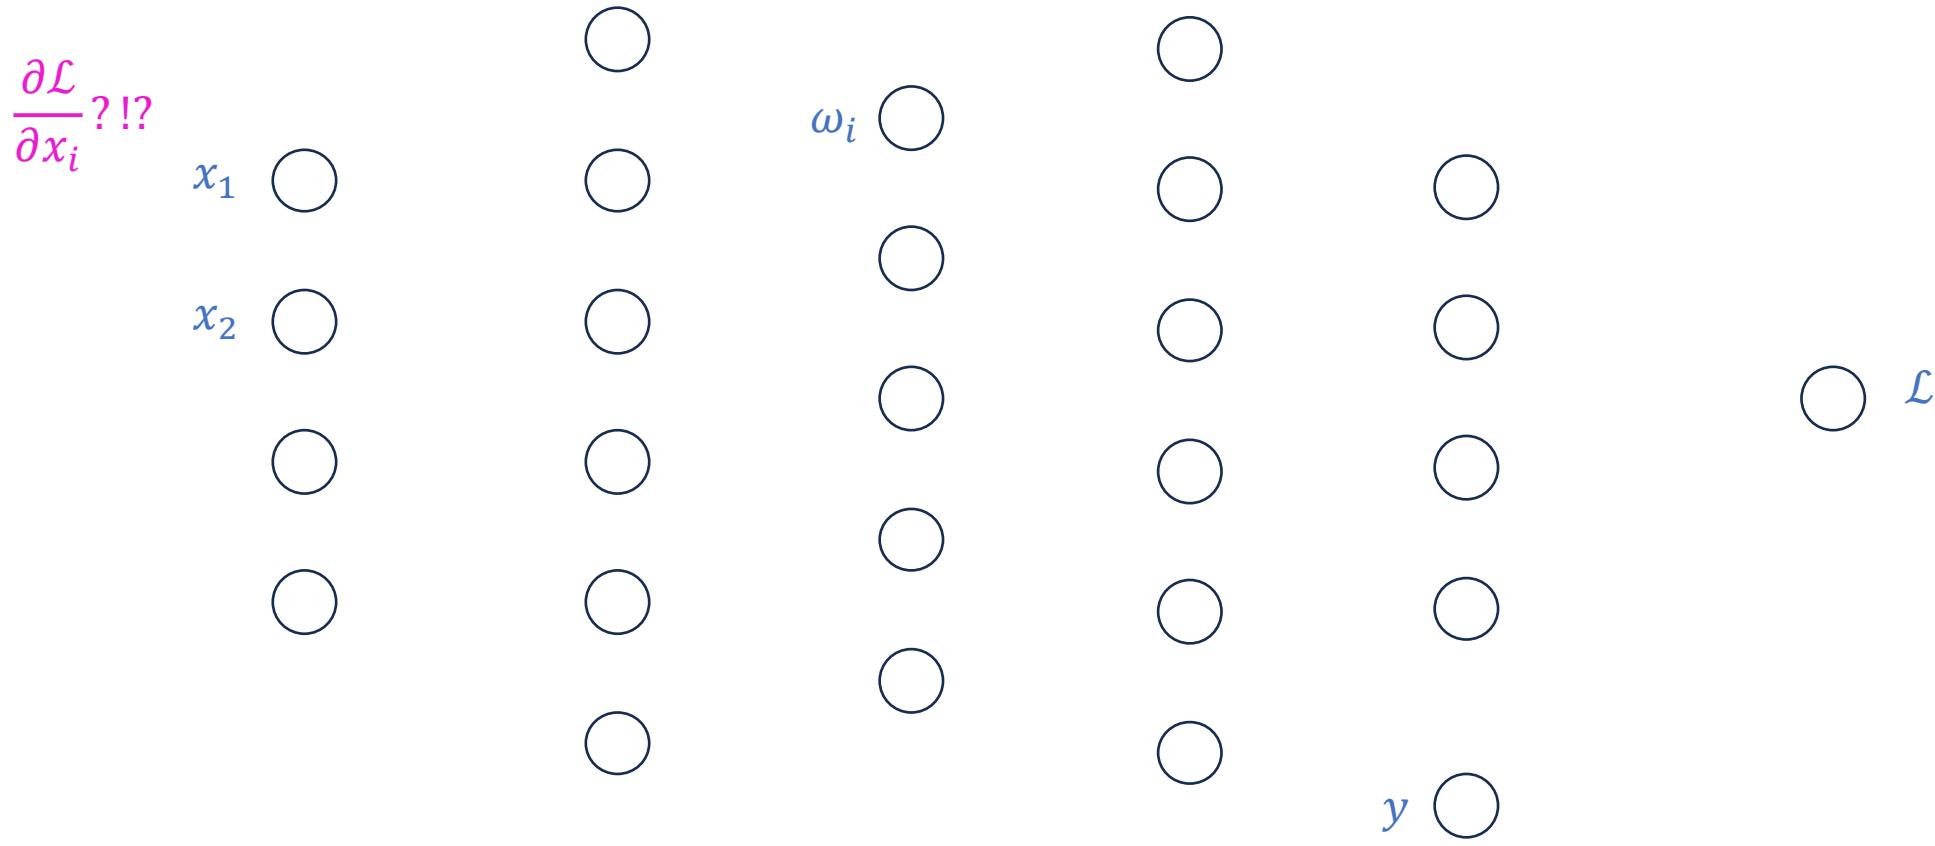
$$\frac{\partial \mathcal{L}}{\partial \omega_i} = \frac{\partial \mathcal{L}}{\partial b} \frac{\partial b}{\partial a} \frac{\partial a}{\partial \omega_i}$$



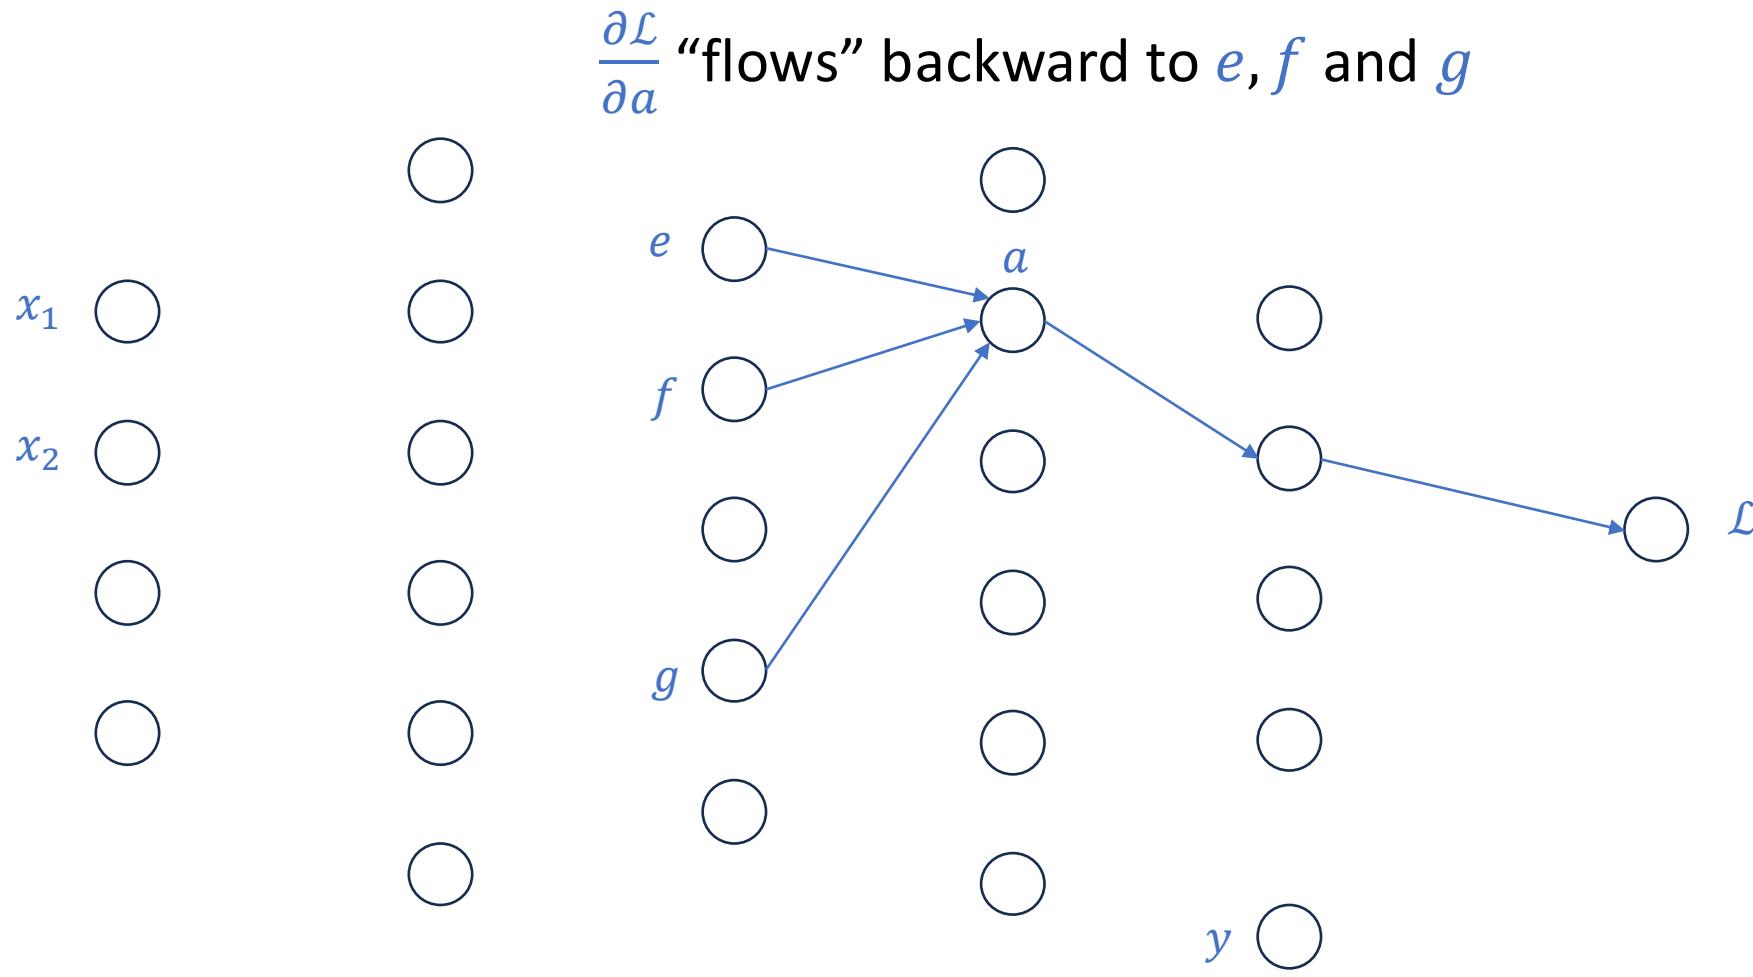
Chain rule



Chain rule

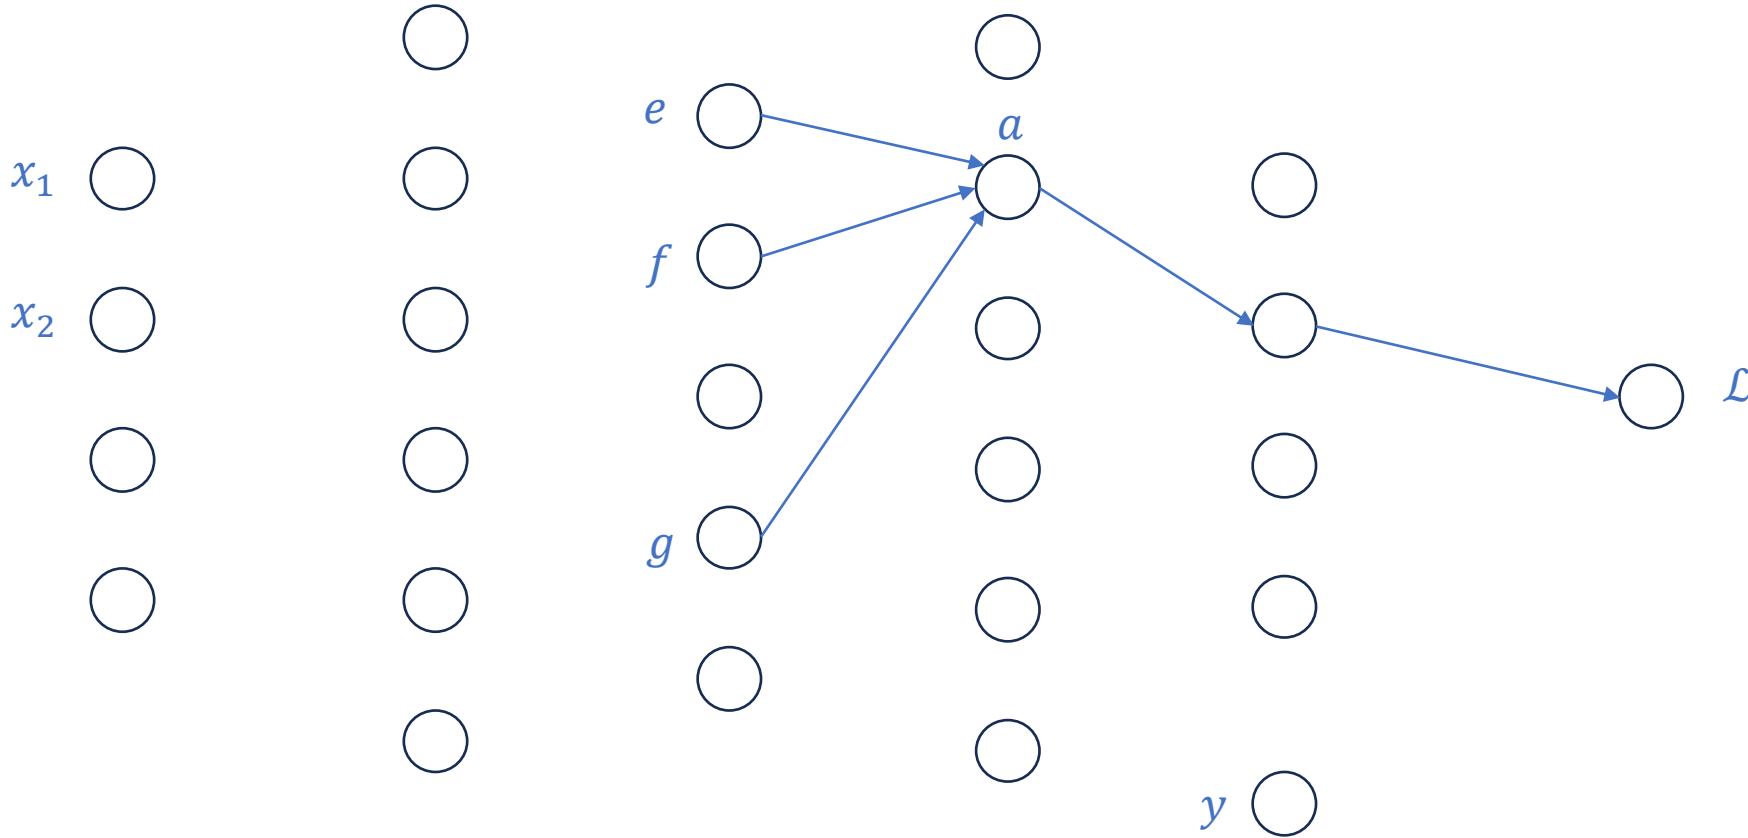


The gradient flow



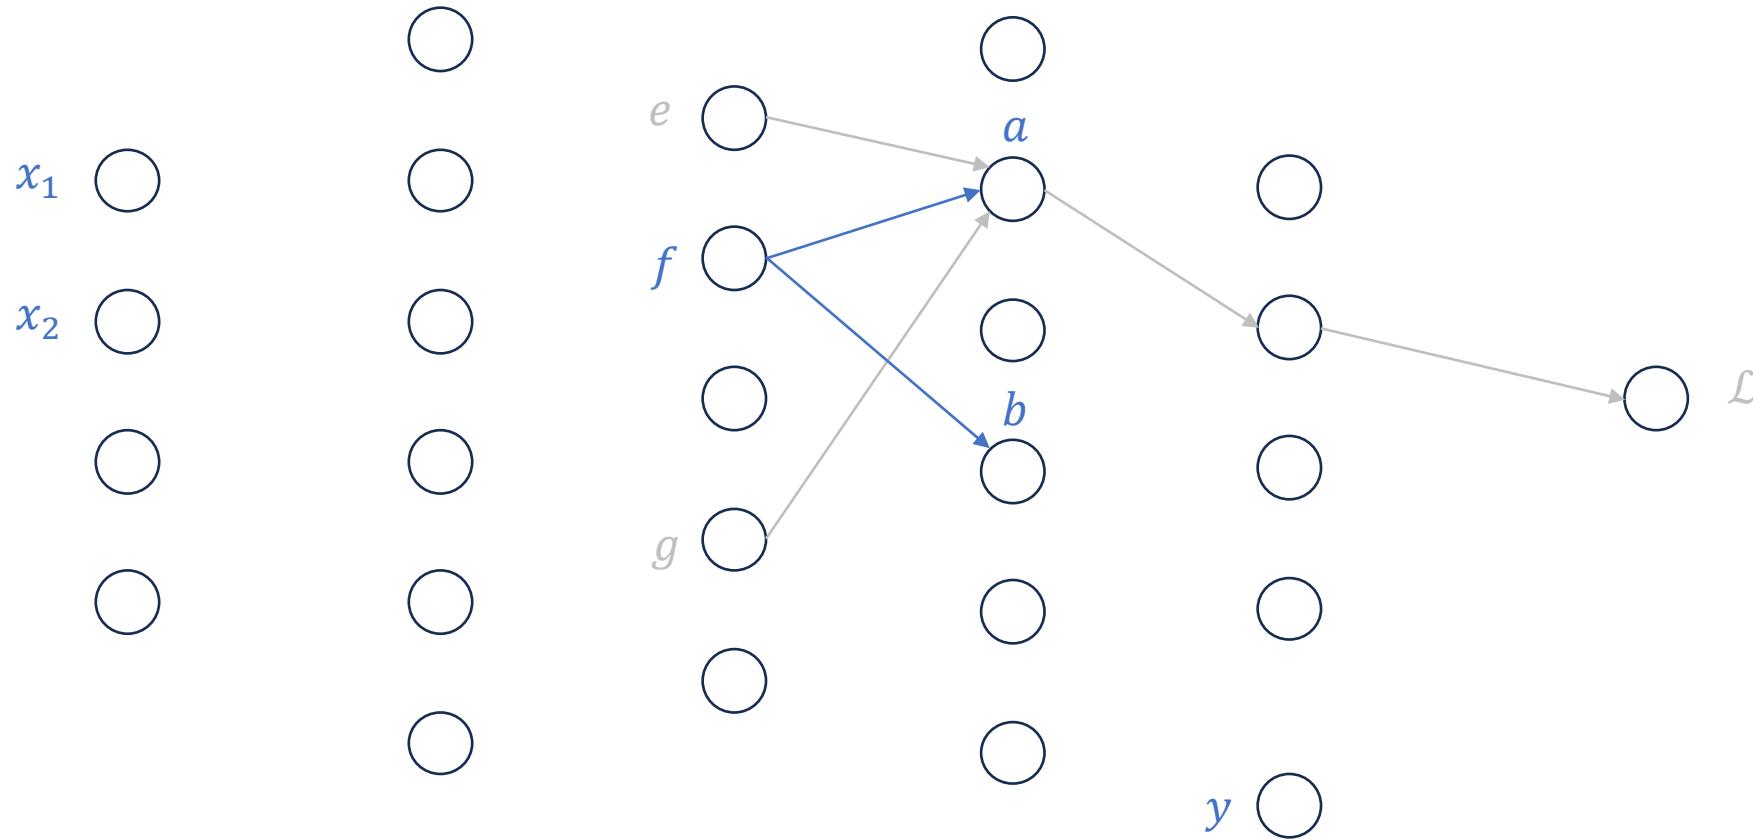
The gradient flow

Given $\frac{\partial \mathcal{L}}{\partial a}$, we can update $\frac{\partial \mathcal{L}}{\partial e}$, $\frac{\partial \mathcal{L}}{\partial f}$ and $\frac{\partial \mathcal{L}}{\partial g}$



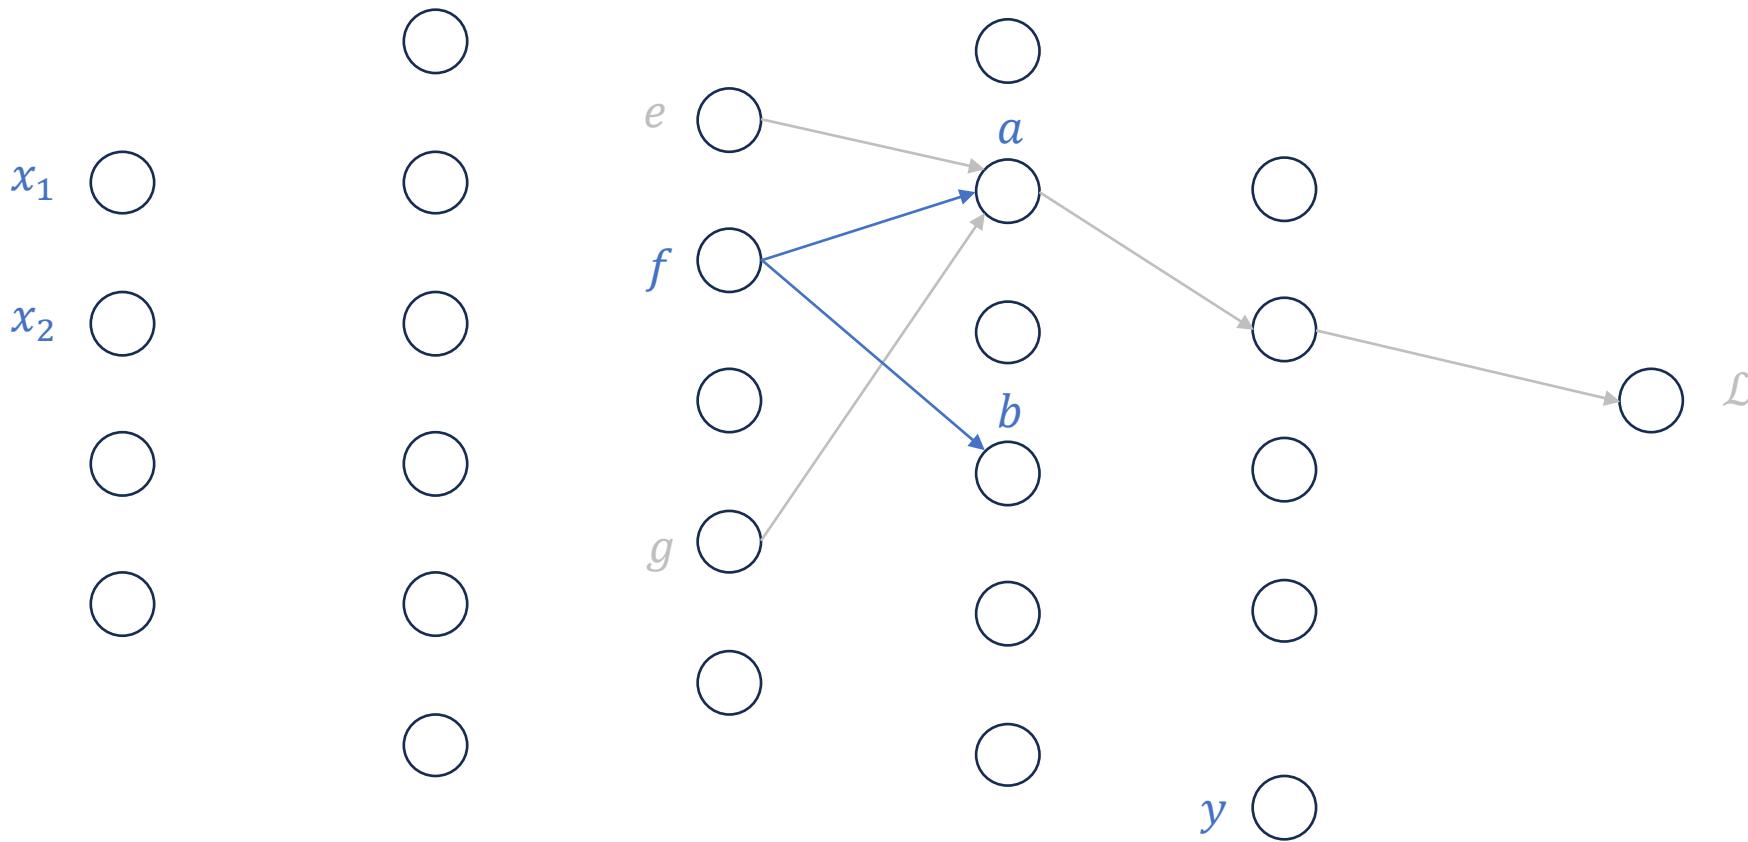
The gradient flow

$\frac{\partial \mathcal{L}}{\partial a}$ and $\frac{\partial \mathcal{L}}{\partial b}$ contribute to $\frac{\partial \mathcal{L}}{\partial f}$



The gradient flow

$$\frac{\partial \mathcal{L}}{\partial f} = \frac{\partial \mathcal{L}}{\partial a} \frac{\partial a}{\partial f} + \frac{\partial \mathcal{L}}{\partial b} \frac{\partial b}{\partial f}$$



Backprop and gradient descent

$$\mathcal{L} = f_{\omega}(x)$$

Assuming f_{ω} is smooth and differentiable w.r.t ω

$$\omega = \omega_0$$

Repeat:

 Compute \mathcal{L}

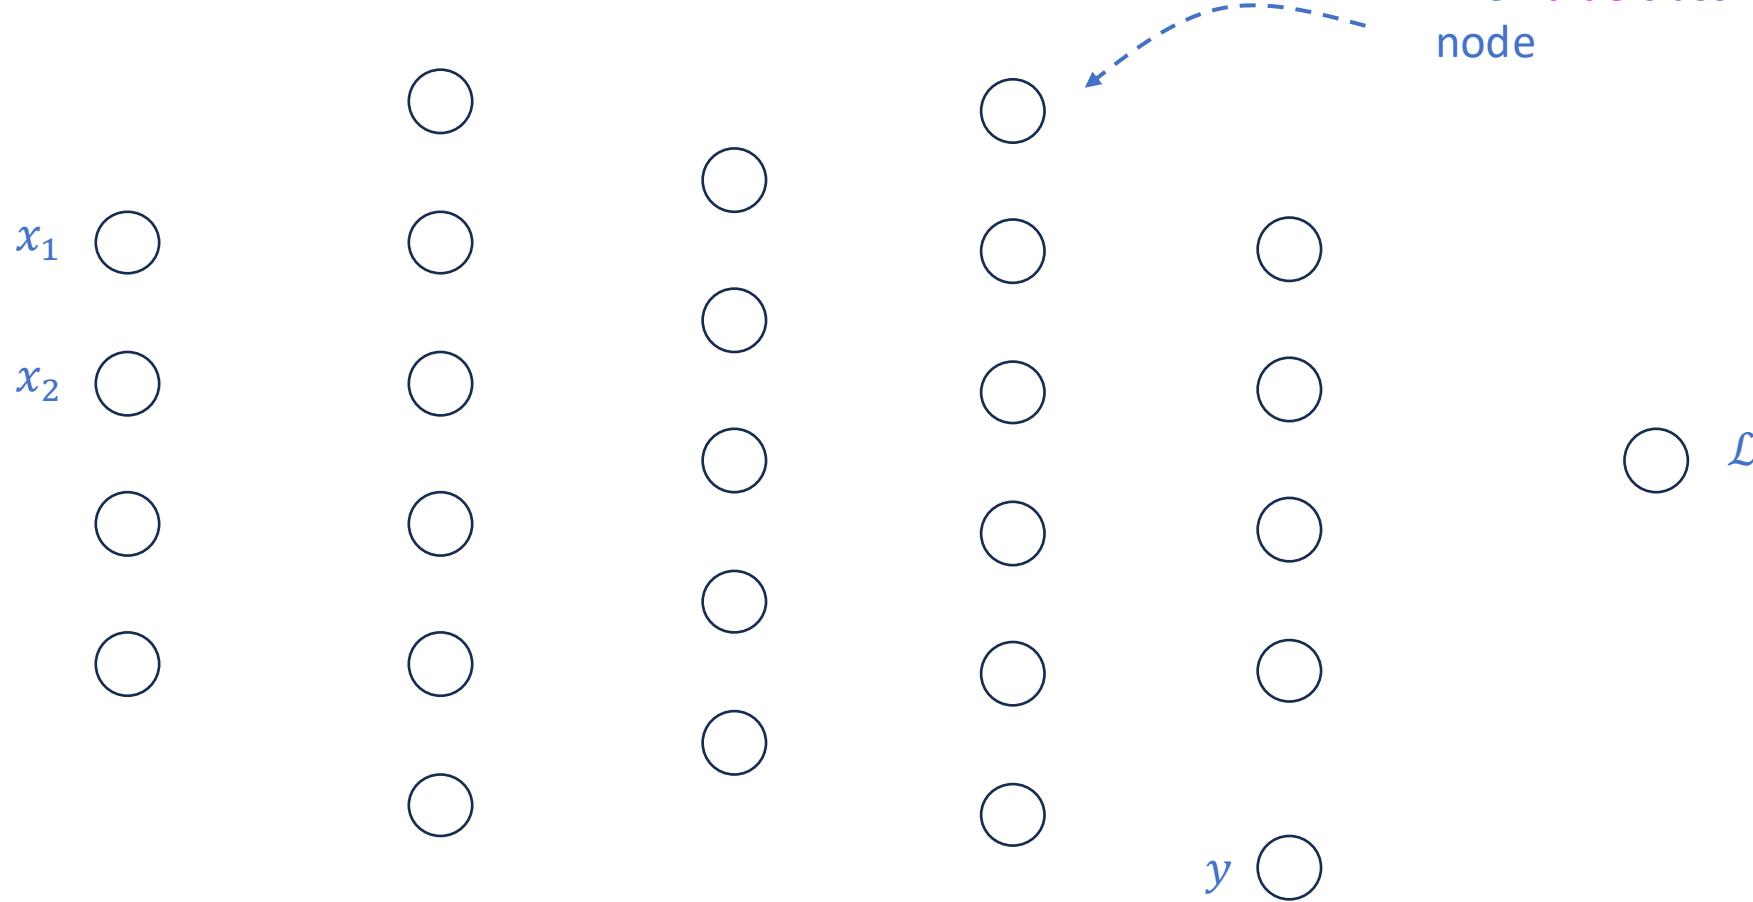
 Reset gradients $\Delta\mathcal{L} = 0$

 Compute $\Delta\mathcal{L}$ (i.e. $\partial\mathcal{L}/\partial\omega_i$)

$\omega_i = \omega_i - \eta \cdot \partial\mathcal{L}/\partial\omega_i$

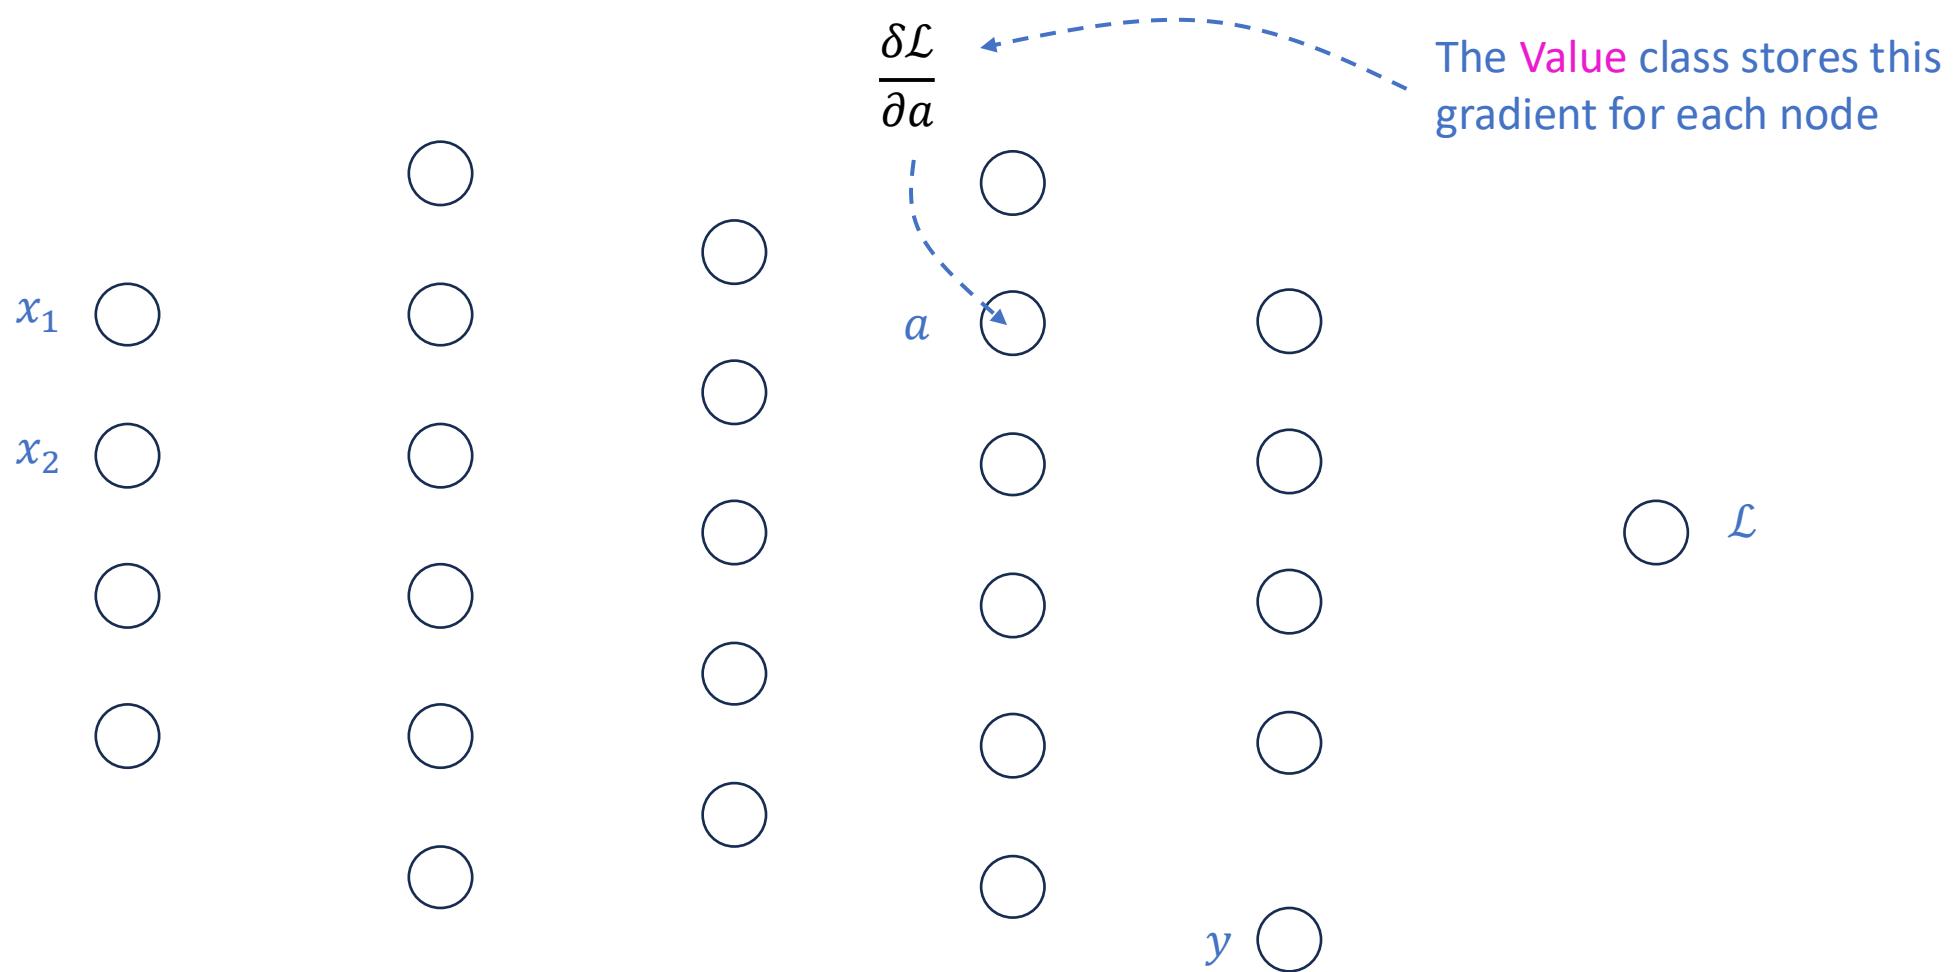
η : learning rate

Let's build it by hand!

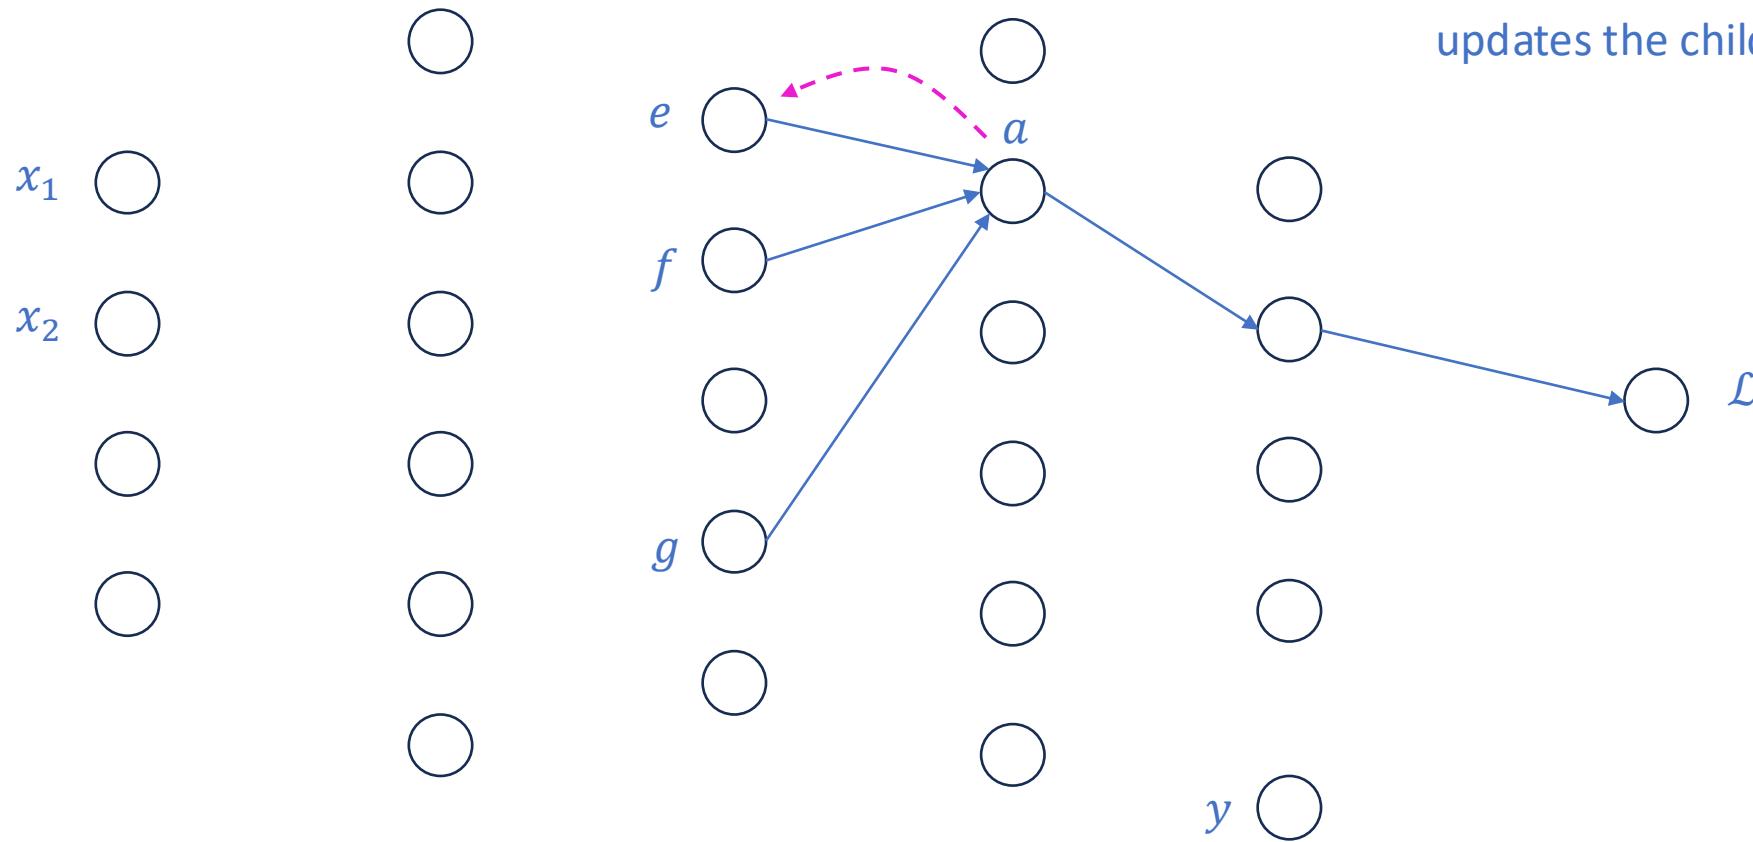


The **Value** class represents a
node

Let's build it by hand!



Let's build it by hand!

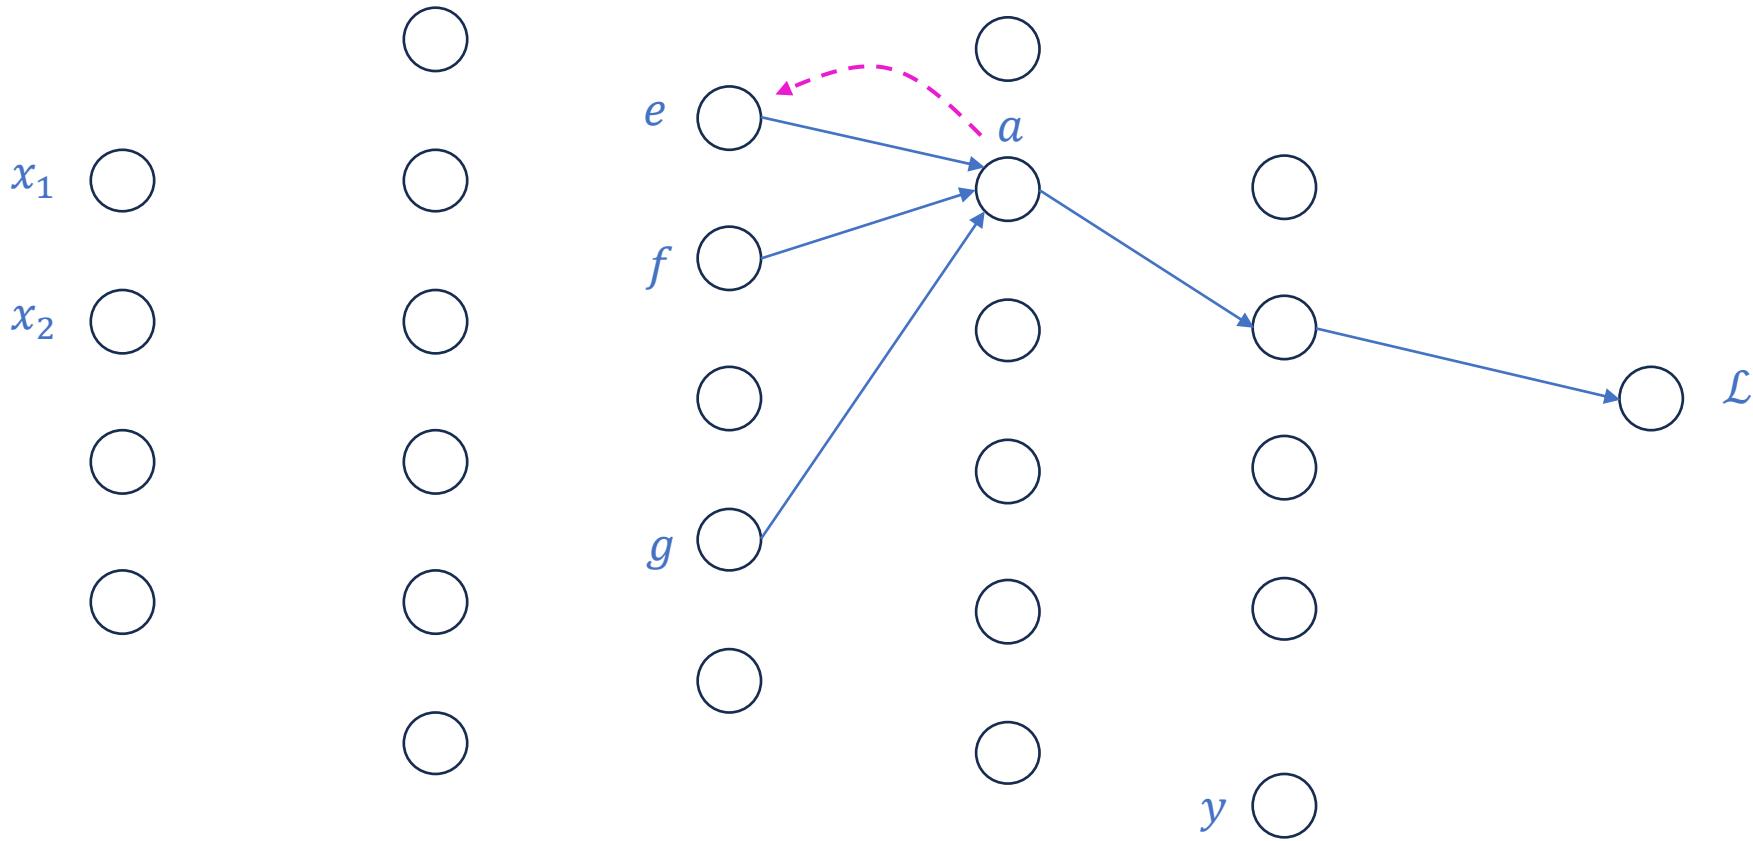


The **Value** class contains a
“backward” function that
updates the children’s gradient

Let's build it by hand!

$$e.\text{grad} += \frac{\partial a}{\partial e} * a.\text{grad}$$

$$\frac{\partial \mathcal{L}}{\partial e} = \frac{\partial \mathcal{L}}{\partial a} \frac{\partial a}{\partial e} + \frac{\partial \mathcal{L}}{\partial b} \frac{\partial b}{\partial e} + \dots$$

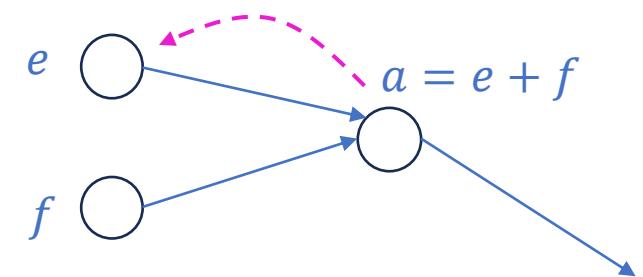


Let's build it by hand!

$$e.\text{grad} += \frac{\partial a}{\partial e} * a.\text{grad}$$

```
class Value:  
    def __init__(self, data):  
        self.data = data  
        self.grad = 0.0
```

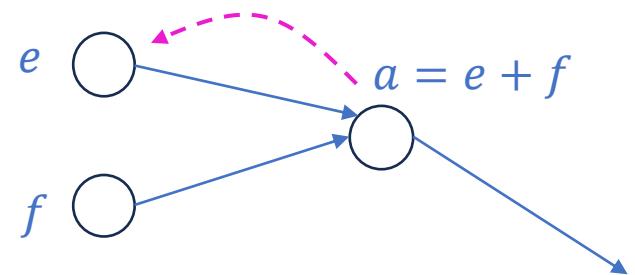
```
    def __add__(self, other):  
        out = Value(self.data + other.data)  
        return out
```



Let's build it by hand!

$$e.\text{grad} += \frac{\partial a}{\partial e} * a.\text{grad}$$

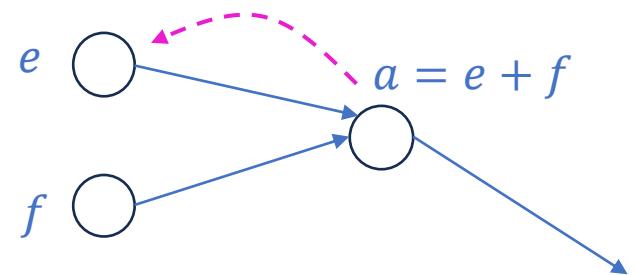
```
class Value:  
    def __init__(self, data):  
        self.data = data  
        self.grad = 0.0  
        self.backward = lambda : None  
  
    def __add__(self, other):  
        out = Value(self.data + other.data)  
        def _backward():  
            self.grad += ???  
            other.grad += ???  
        out.backward = _backward  
        return out
```



Let's build it by hand!

$$e.\text{grad} += \frac{\partial a}{\partial e} * a.\text{grad}$$

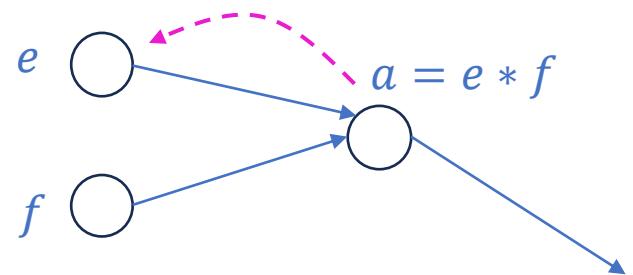
```
class Value:  
    def __init__(self, data):  
        self.data = data  
        self.grad = 0.0  
        self.backward = lambda : None  
  
    def __add__(self, other):  
        out = Value(self.data + other.data)  
        def _backward():  
            self.grad += 1.0 * out.grad  
            other.grad += 1.0 * out.grad  
        out.backward = _backward  
        return out
```



Let's build it by hand!

$$e.\text{grad} += \frac{\partial a}{\partial e} * a.\text{grad}$$

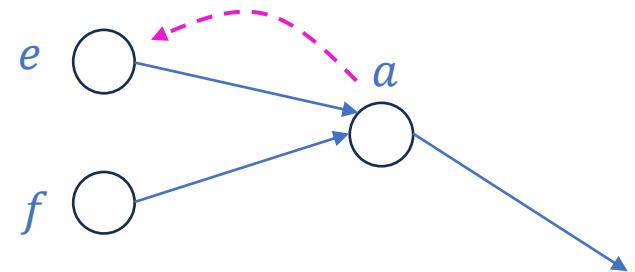
```
class Value:  
    def __init__(self, data):  
        self.data = data  
        self.grad = 0.0  
        self.backward = lambda : None  
  
    def __mul__(self, other):  
        out = Value(self.data * other.data)  
        def _backward():  
            self.grad += ??? * out.grad  
            other.grad += ??? * out.grad  
        out.backward = _backward  
        return out
```



Let's build it by hand!

$$e.\text{grad} += \frac{\partial a}{\partial e} * a.\text{grad}$$

```
class Value:  
    def __init__(self, data):  
        self.data = data  
        self.grad = 0.0  
        self.backward = lambda : None  
  
    def __mul__(self, other):  
        out = Value(self.data * other.data)  
        def _backward():  
            self.grad += other.data * out.grad  
            other.grad += self.data * out.grad  
        out.backward = _backward  
        return out
```

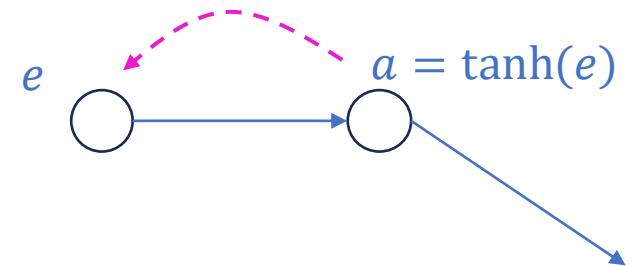


$$a = e * f$$

$$\frac{\partial a}{\partial e} = f$$

Let's build it by hand!

$$e.\text{grad} += \frac{\partial a}{\partial e} * a.\text{grad}$$



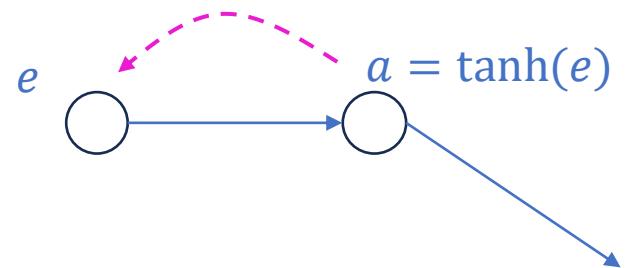
```
class Value:  
    def __init__(self, data):  
        self.data = data  
        self.grad = 0.0  
        self.backward = lambda : None  
  
    def __tanh__(self):  
        out = Value(???)  
        def _backward():  
            self.grad += ??? * out.grad  
        out.backward = _backward  
        return out
```

$$a = \tanh(e)$$

$$\frac{\partial a}{\partial e} = ???$$

Let's build it by hand!

$$e.\text{grad} += \frac{\partial a}{\partial e} * a.\text{grad}$$

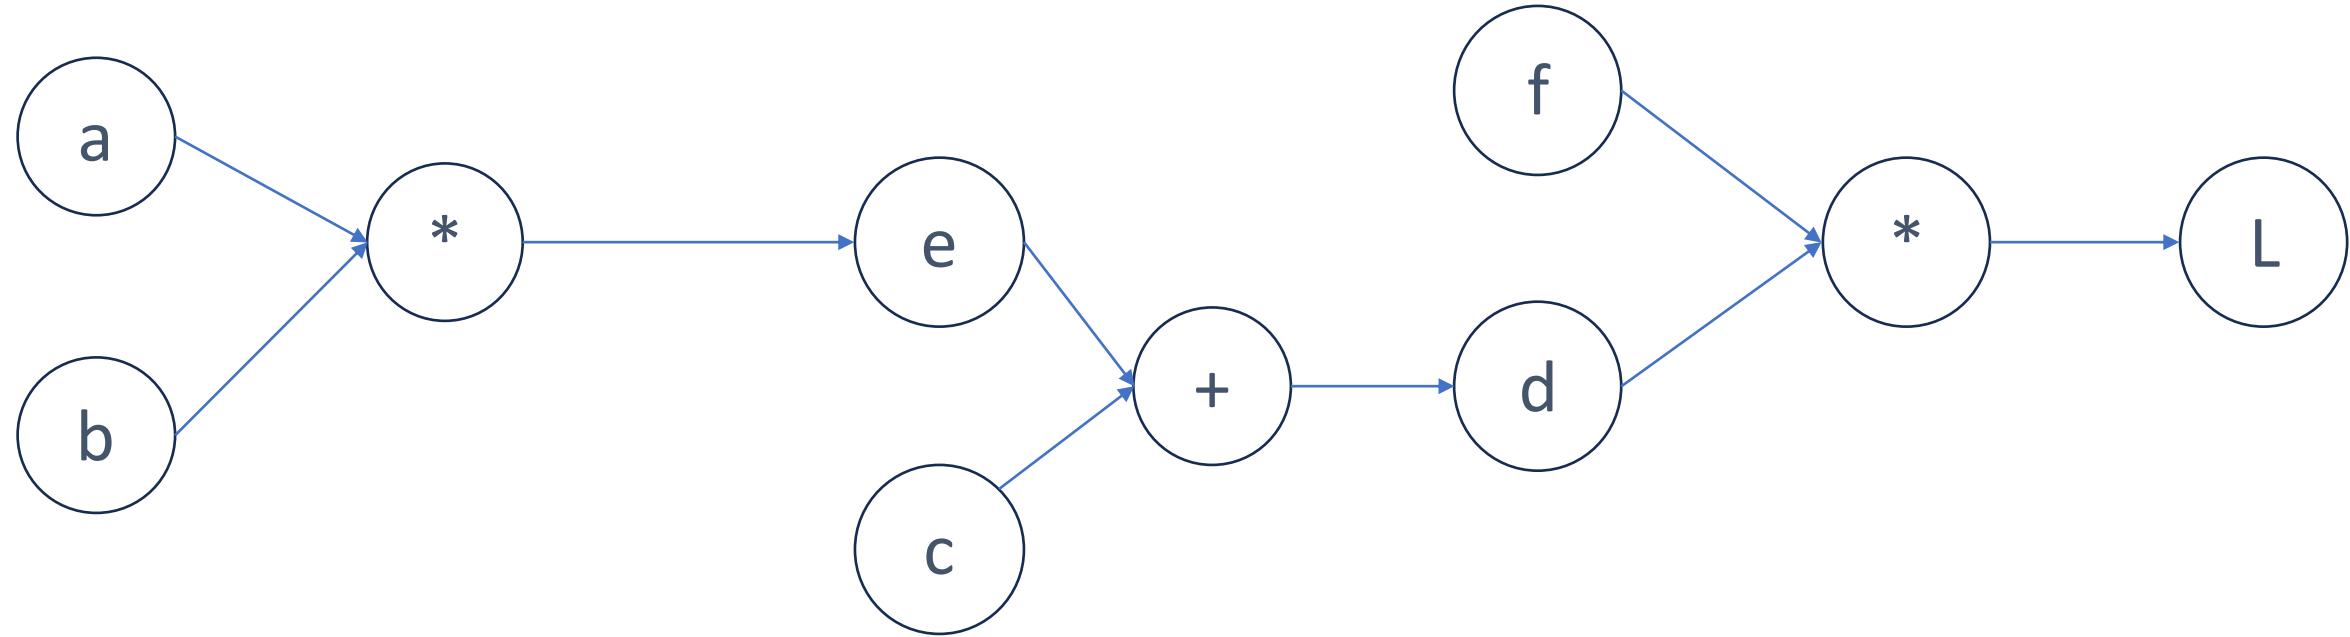


```
class Value:  
    def __init__(self, data):  
        self.data = data  
        self.grad = 0.0  
        self.backward = lambda : None  
  
    def __tanh__(self):  
        t = (math.exp(2*x) - 1) / (math.exp(2*x) + 1)  
        out = Value(t, (self, ), 'tanh')  
        def __backward__():  
            self.grad += (1.0 - t**2) * out.grad  
        out.backward = __backward__  
        return out
```

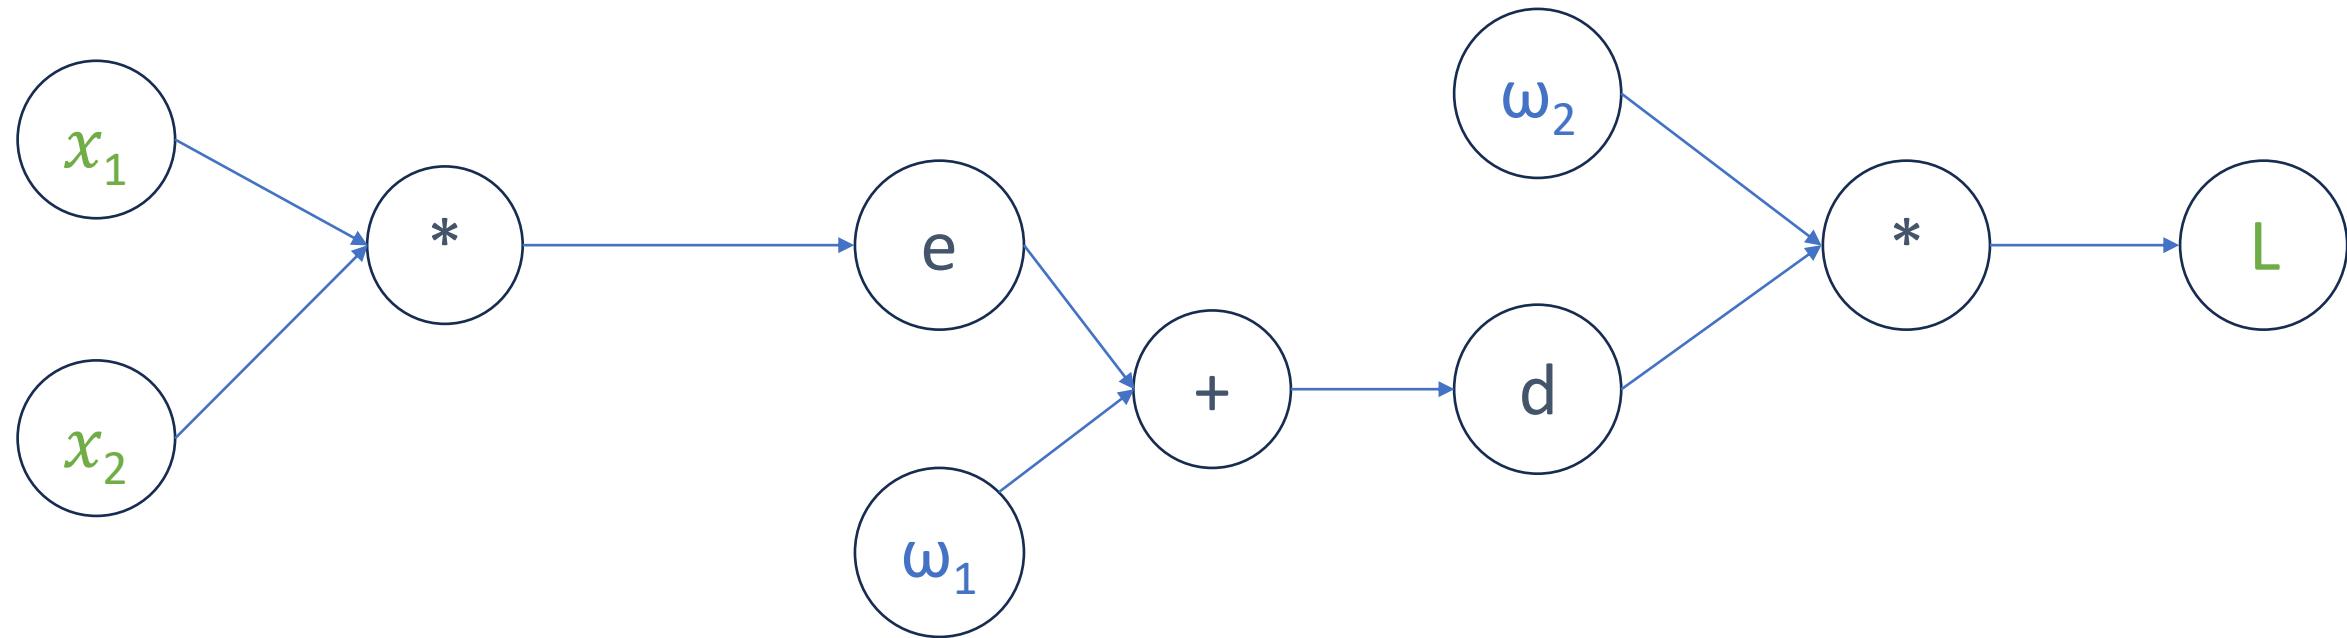
$$a = \tanh(e)$$

$$\frac{\partial a}{\partial e} = 1 - a^2$$

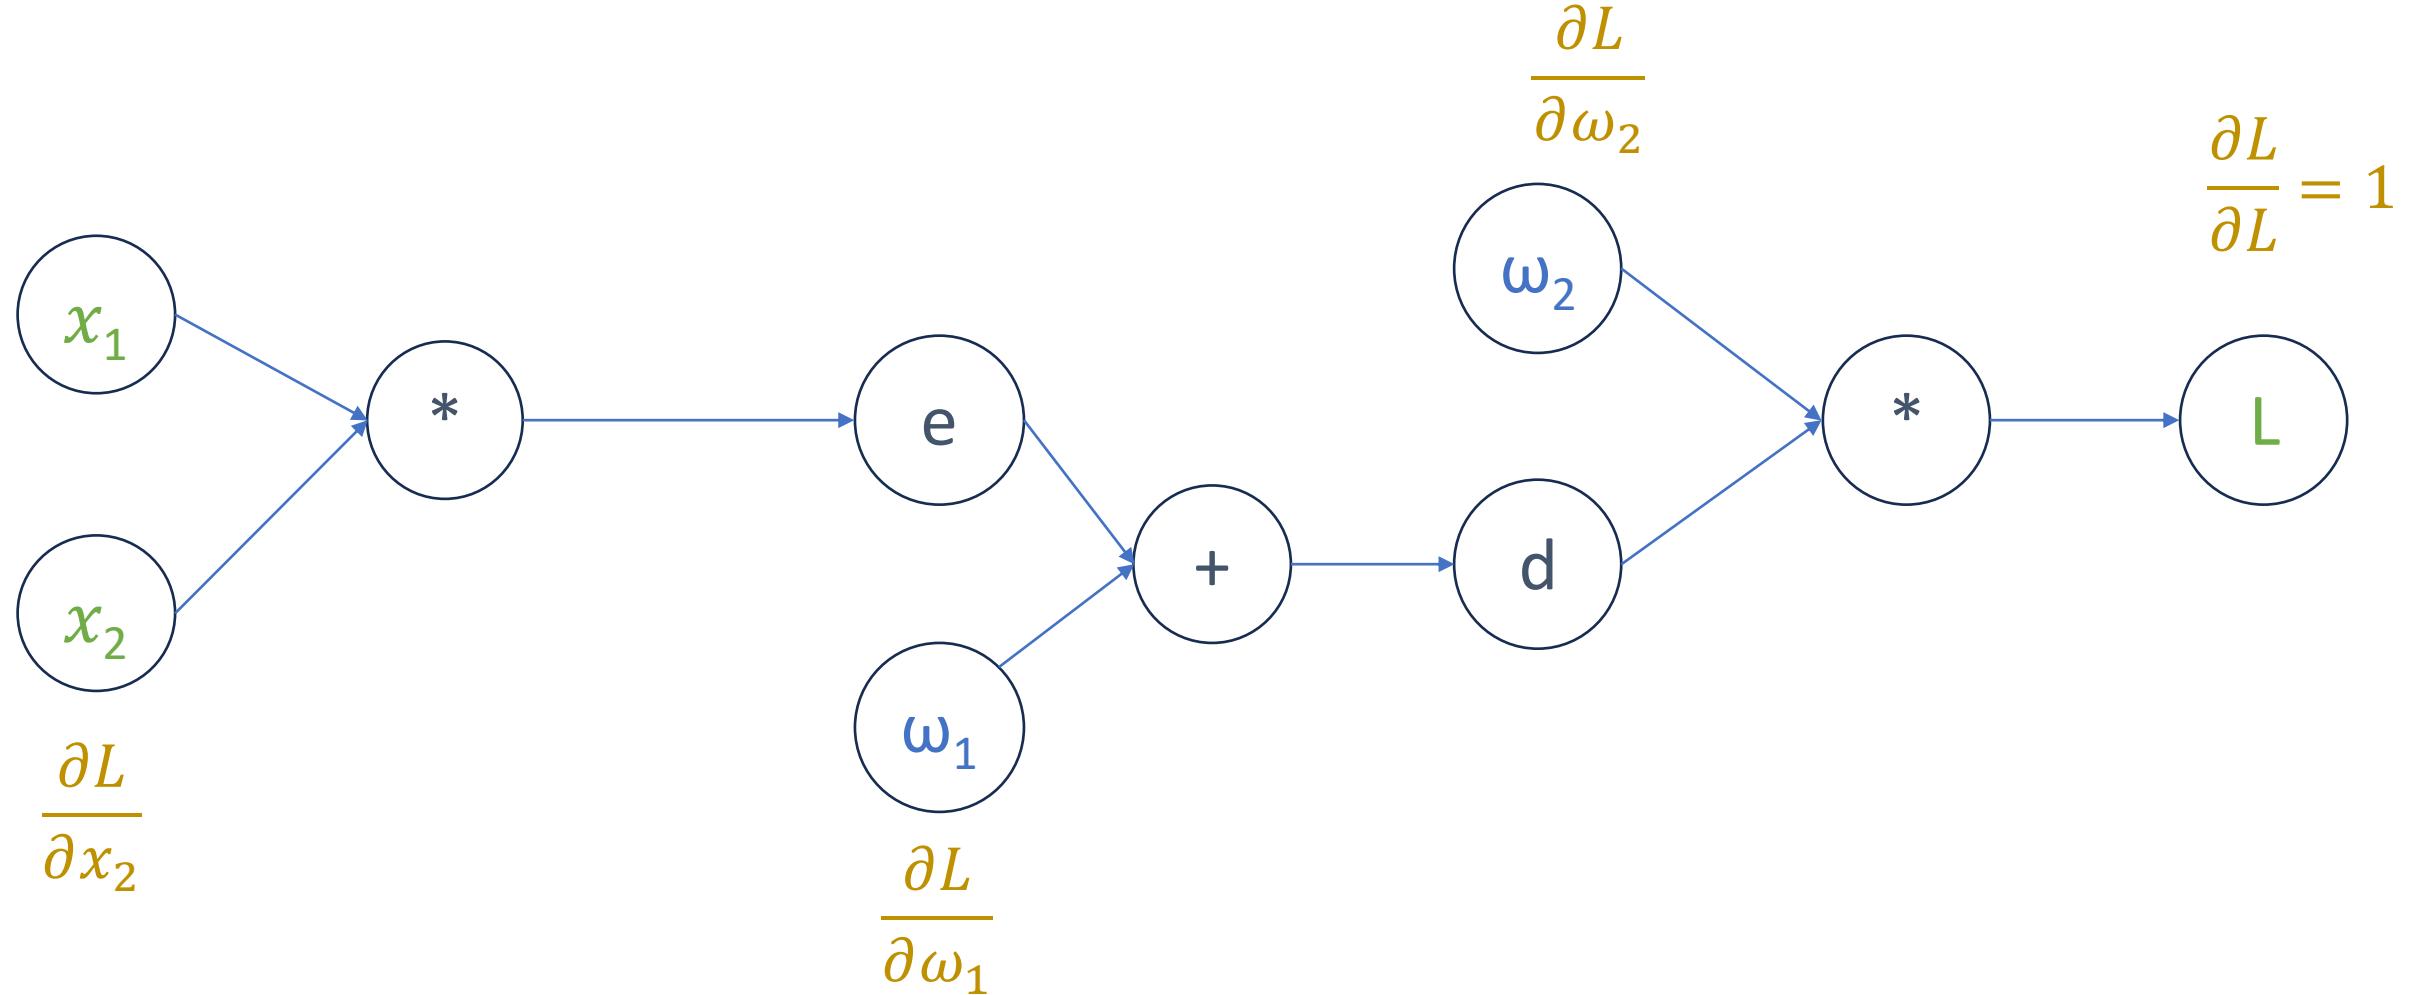
The `lol()` function



The `lol()` function



The `lol()` function

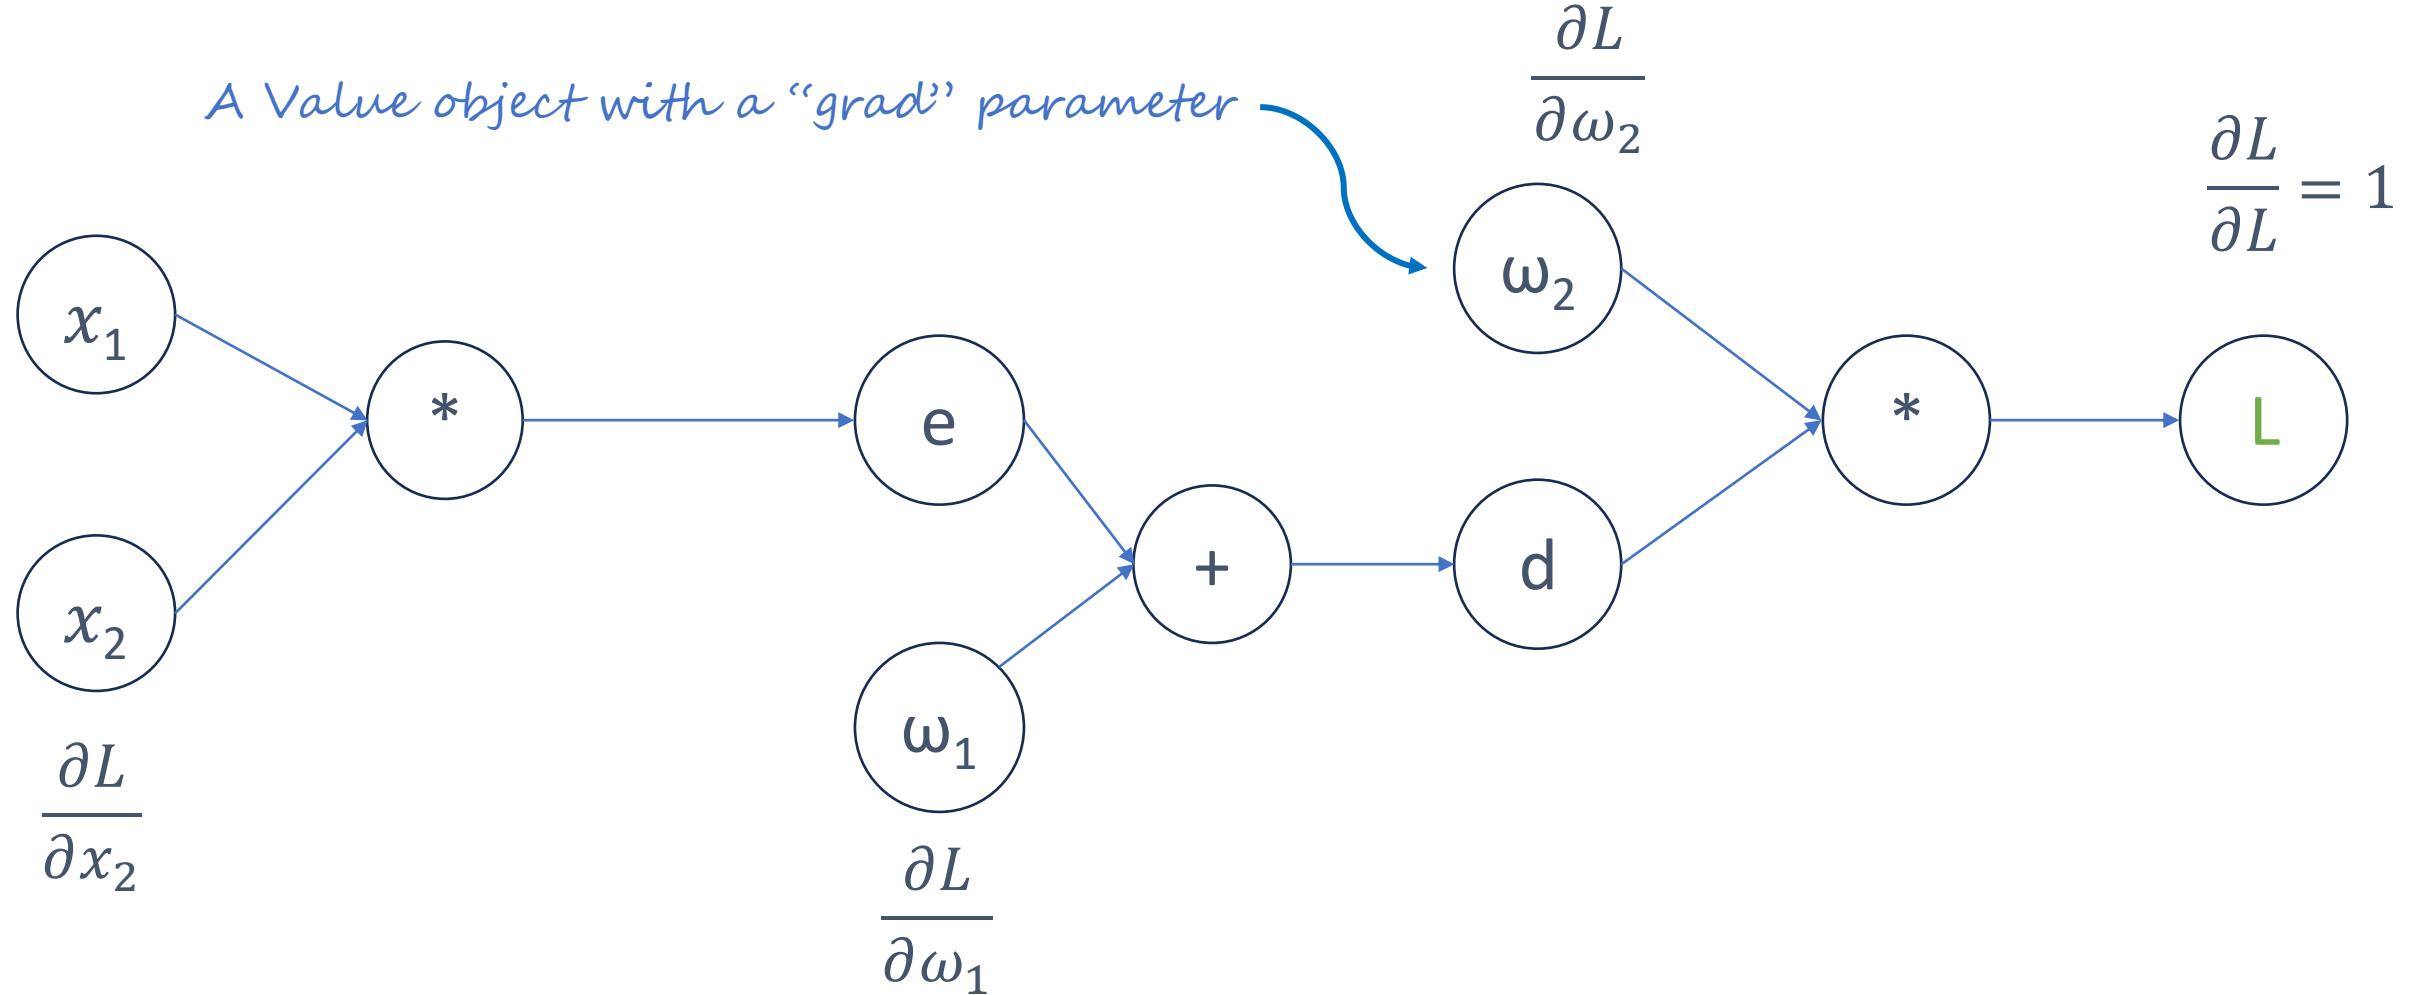


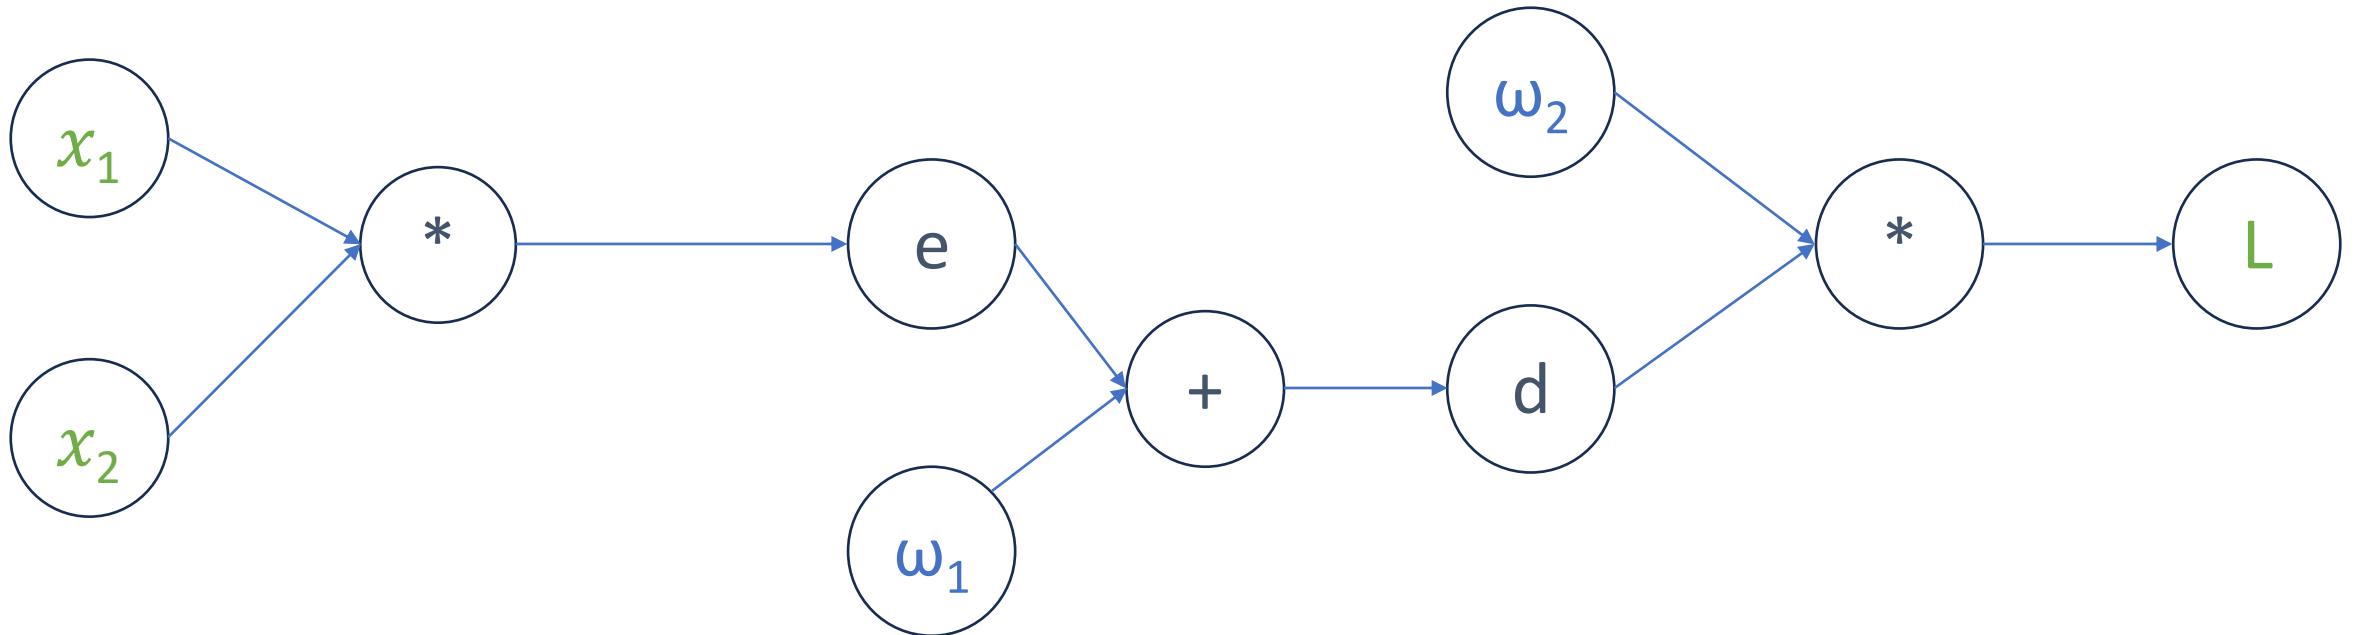
Gradient descent for the `lol()` function

$$\omega_1 = \omega_1 - \eta \frac{\partial L}{\partial \omega_1}$$

$$\omega_2 = \omega_2 - \eta \frac{\partial L}{\partial \omega_2}$$

The lol() function





```
class Neuron:

    def __init__(self, nin):
        self.w = [Value(random.uniform(-1,1)) for _ in range(nin)]
        self.b = Value(random.uniform(-1,1))

    def __call__(self, x):
        # w * x + b
        act = ??? # Placeholder for activation function
        out = act.tanh()
        return out

    def parameters(self):
        return self.w + [self.b]
```

Mean square error loss

$$\ell = \sum (y - \hat{y})^2$$

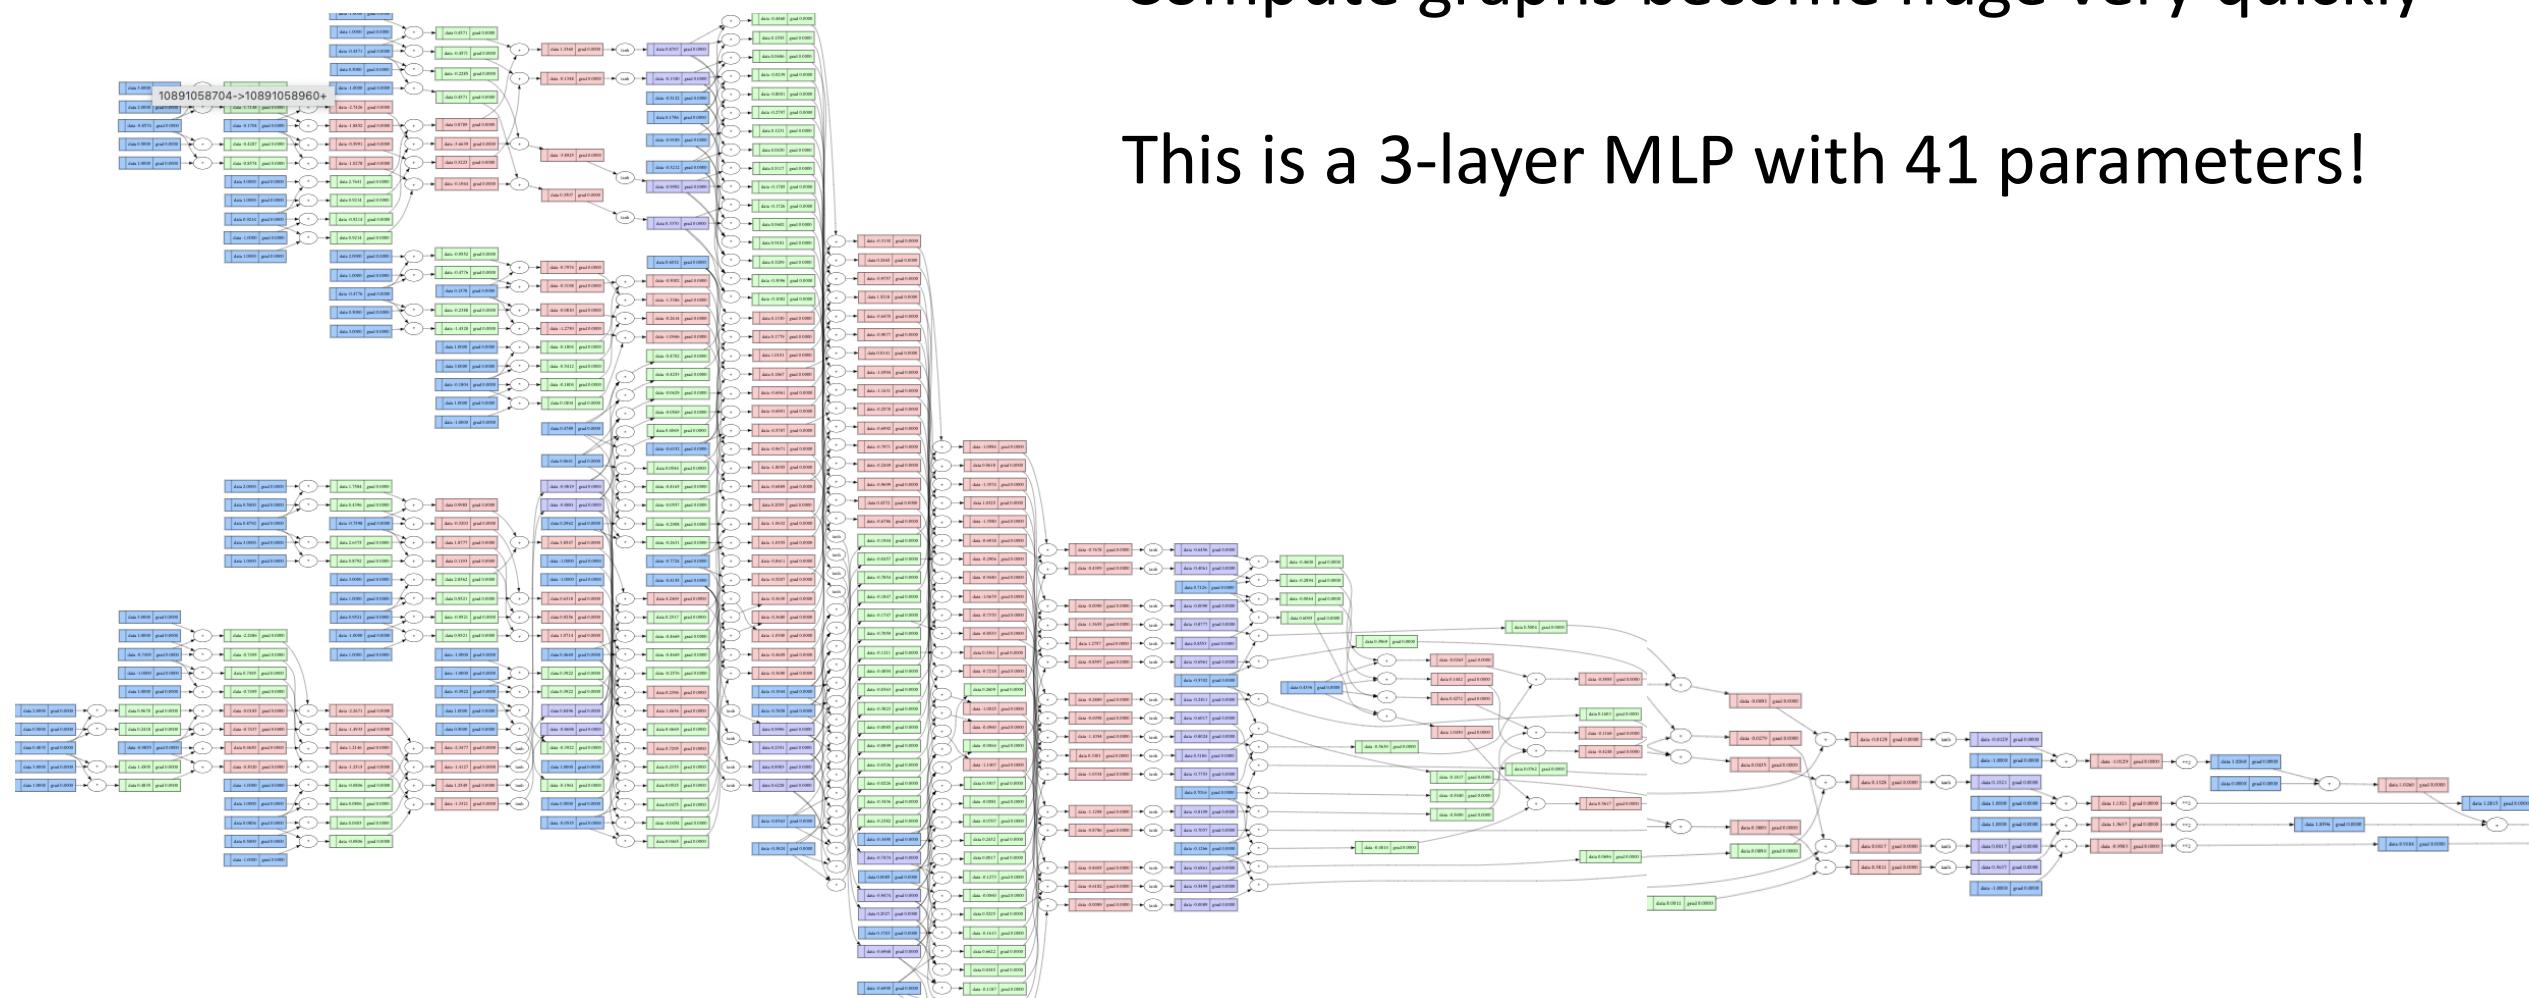
```
n = MLP(3, [4, 4, 1])
```

```
ypred = [n(x) for x in xs]
```

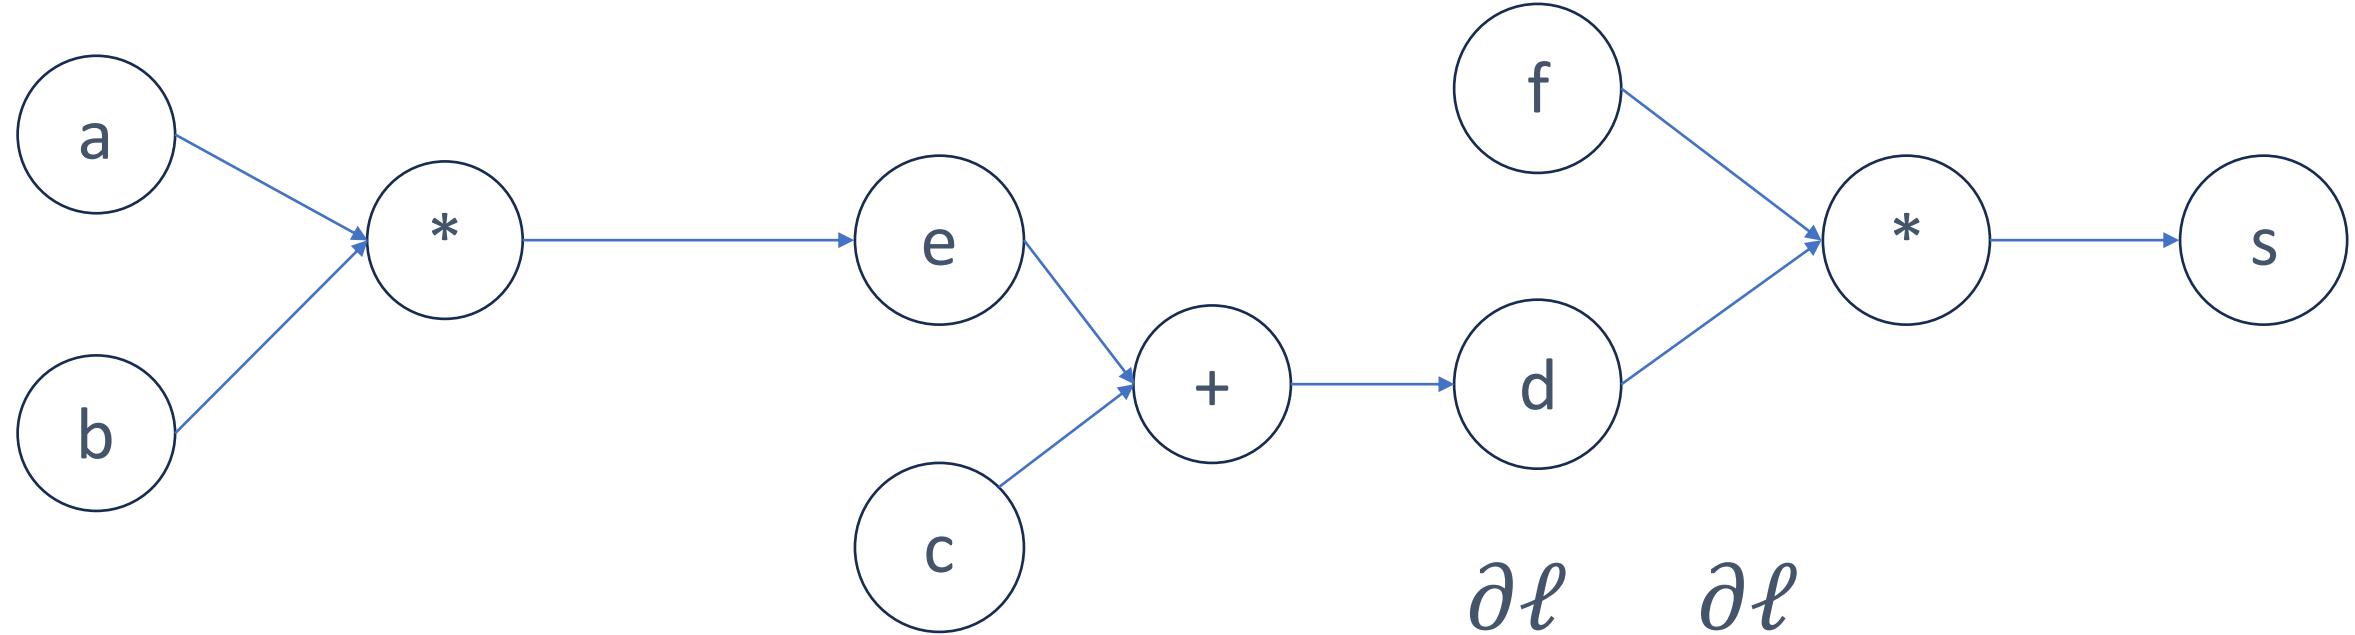
```
loss = sum([(a-b)**2 for (a,b) in zip(ypred, ys)])
```

Compute graphs become huge very quickly

This is a 3-layer MLP with 41 parameters!



The forward pass is critical to update the values in the network



$$\frac{\partial \ell}{\partial f} = \frac{\partial \ell}{\partial s} * d$$

Stochastic gradient descent

Iterate over **batches** of samples instead of the whole dataset at once

1. Scales to large datasets (that don't fit in memory)
2. Adds randomness to the process
3. Adds implicit regularization

Conclusion

Neural networks are **compute graphs**.

Gradient descent minimizes a loss function over the network's parameters.

Back-propagation allows **efficient learning** (tuning of the network).

Let's practice with a simple Multi-Layer Perception (MLP).

Universal approximation theorem

For any continuous function, there exists a shallow network that can approximate this function to any specified precision.

Universal approximation theorem

Universal approximation theorem (Uniform non-affine activation, arbitrary depth, constrained width). Let \mathcal{X} be a compact subset of \mathbb{R}^d . Let $\sigma : \mathbb{R} \rightarrow \mathbb{R}$ be any non-affine continuous function which is continuously differentiable at at least one point, with nonzero derivative at that point. Let $\mathcal{N}_{d,D:d+D+2}^\sigma$ denote the space of feed-forward neural networks with d input neurons, D output neurons, and an arbitrary number of hidden layers each with $d + D + 2$ neurons, such that every hidden neuron has activation function σ and every output neuron has the identity as its activation function, with input layer ϕ and output layer ρ . Then given any $\varepsilon > 0$ and any $f \in C(\mathcal{X}, \mathbb{R}^D)$, there exists $\hat{f} \in \mathcal{N}_{d,D:d+D+2}^\sigma$ such that

$$\sup_{x \in \mathcal{X}} \|\hat{f}(x) - f(x)\| < \varepsilon.$$

In other words, \mathcal{N} is dense in $C(\mathcal{X}; \mathbb{R}^D)$ with respect to the topology of uniform convergence.

No free-lunch theorem

Every learning algorithm is as good as any other when averaged over all sets of problems.

You can't just learn « purely from data » without bias.

Wolpert, D. H.; Macready, W. G. (1997). ["No Free Lunch Theorems for Optimization"](#). *IEEE Transactions on Evolutionary Computation*. 1: 67–82.